
NVIDIA
INFINIBAND HPC CLOUD SOLUTION

Aganda:

1. Infiniband如何在公有云建设，如何部署
2. Infiniband建设/部署的拓扑，最小启动建设单元
3. Infiniband如何监控运维
4. IB部署方案分享

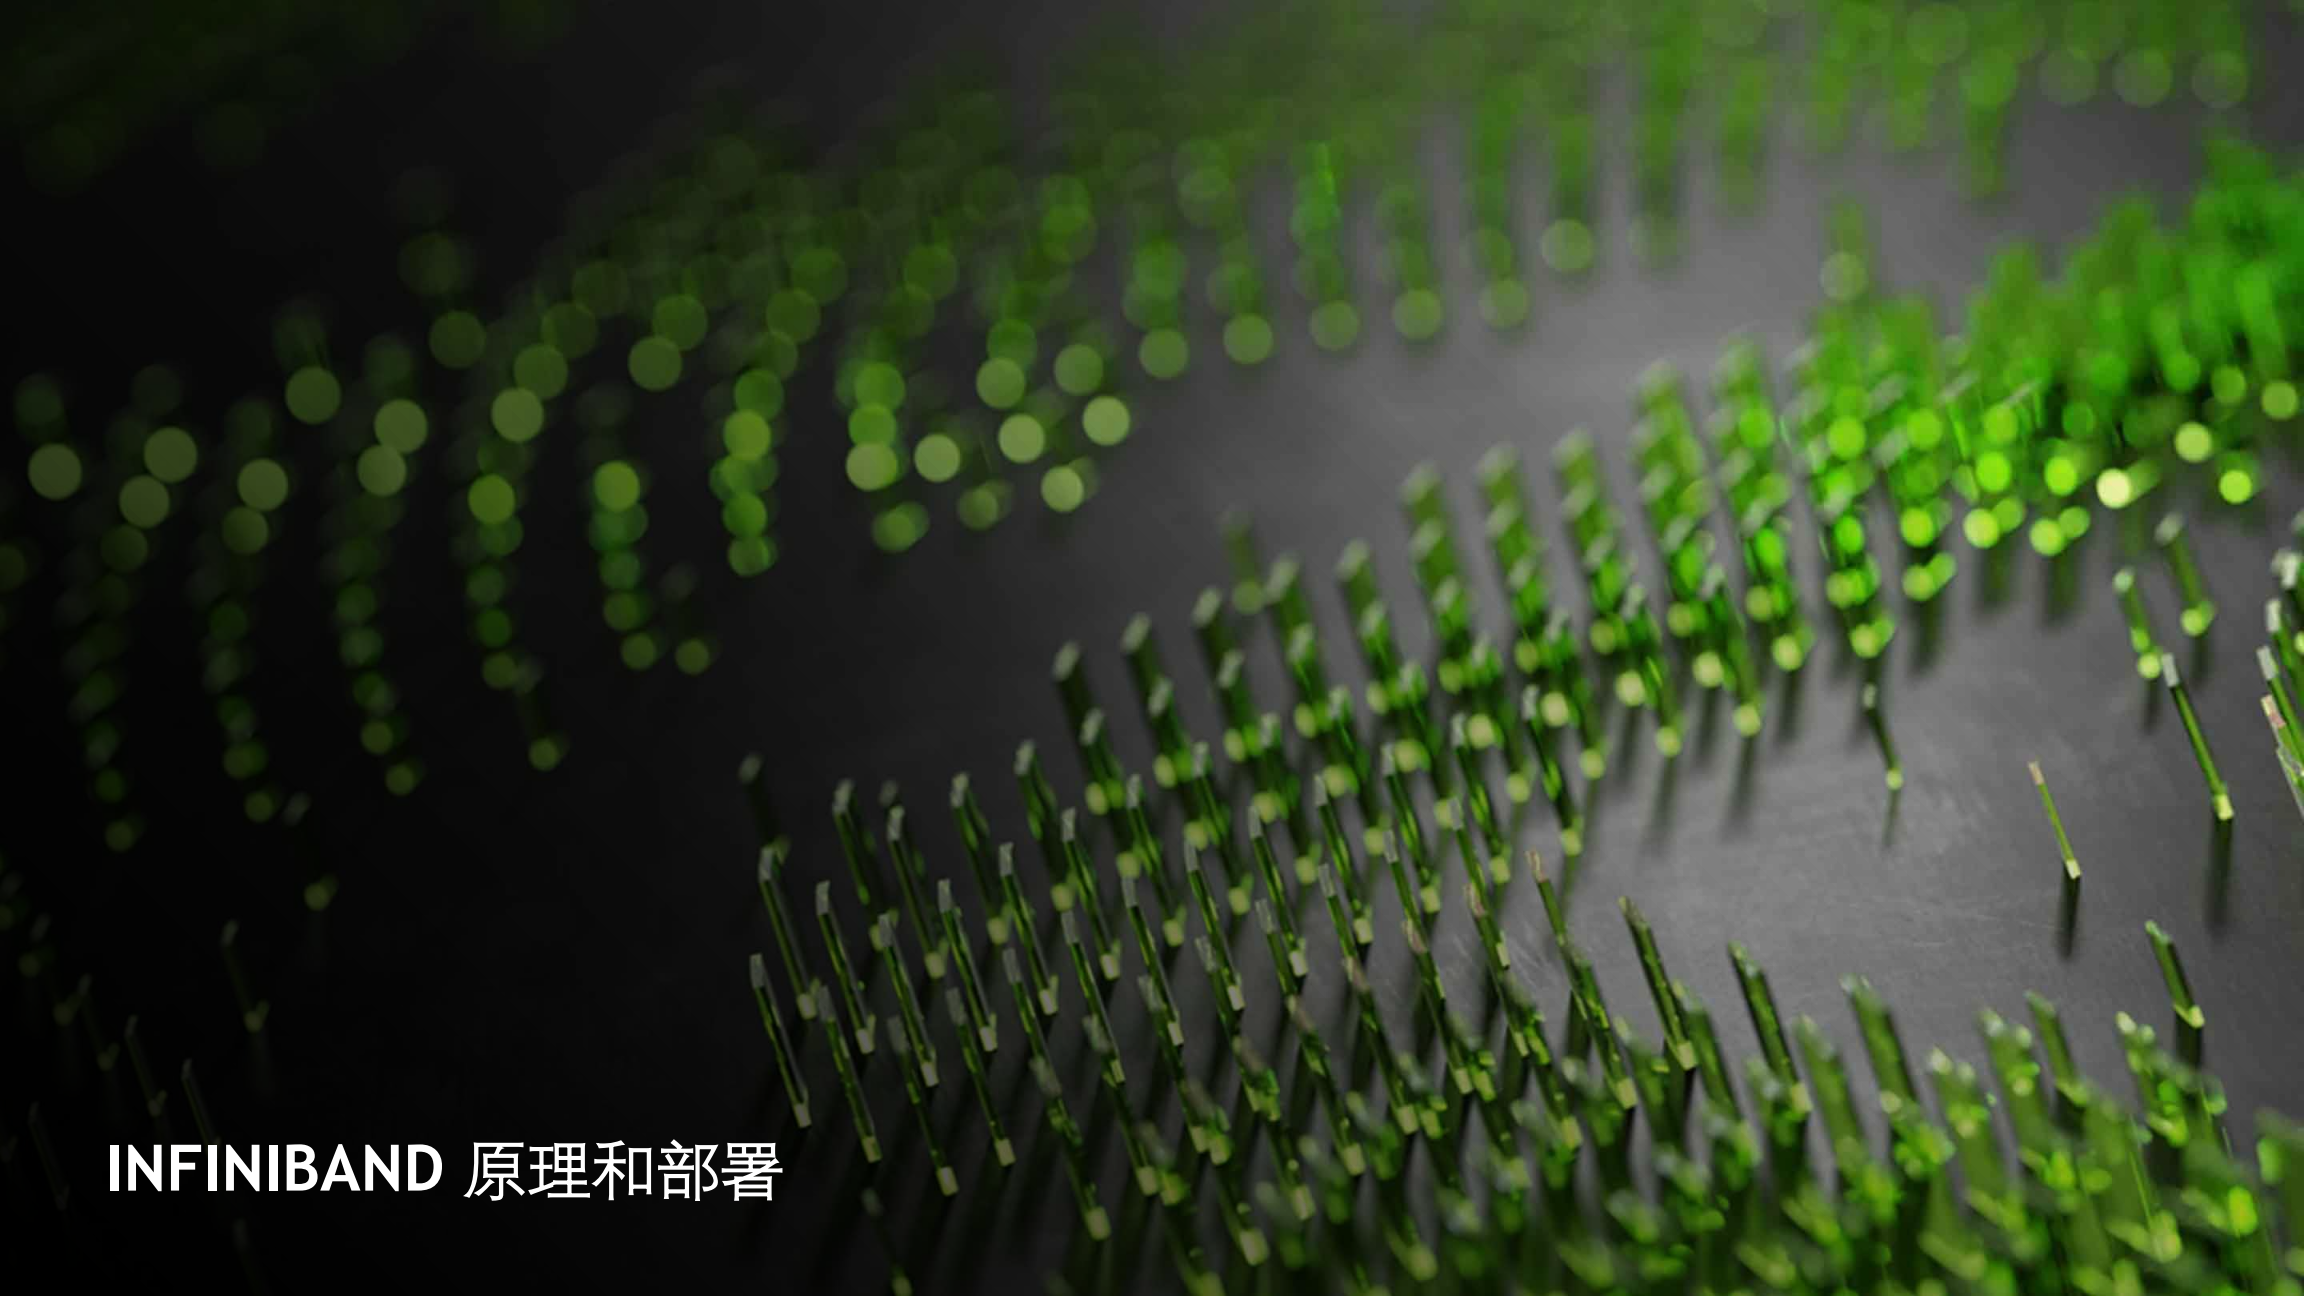

INFINIBAND 原理和部署

SUBNET MODEL

- Subnet = HCAs and interconnected through switches
- Each subnet has its own LID space
- Each subnet has at least one SM and exactly one (logical) Master SM

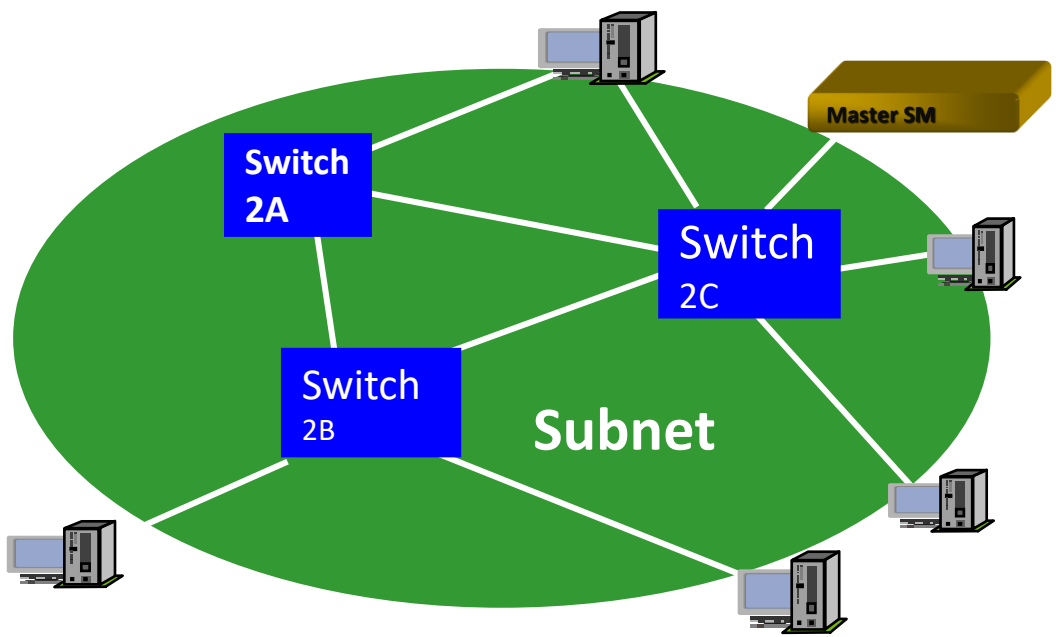

INFINIBAND SUBNET MANAGER

- SM requests like devices responses will include :
  - Node Info
  - Ports info

1. SM – I am requesting info via port number 1
2. Switch A – I am responding via port number 3 , I have an Active port Number 4
3. SM – I am requesting info via :
   - my port 1- next switch (A) port 4
4. SWITCH B – I am a switch responding via my port 3 via next switch (A) port 3
   - I have a live port port 4
5. SM – I am requesting info via :
   - my port 1- next switch (A) port 4 ,next switch (B) port 4
6. Host – I am a CA , responding via my port 1 , next switch (B) port 3 , next switch (A) port 3

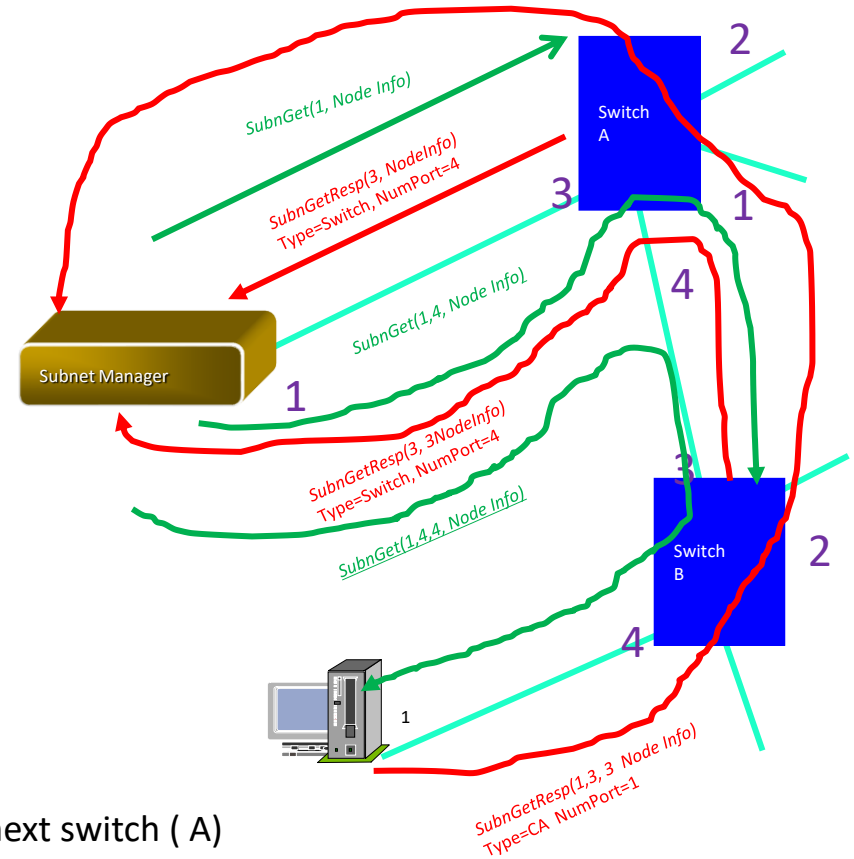

IB FORWARDING TABLE

- After the SM finished gathering
  all Fabric information , including direct route tables ,
  it assigns a LID to each one of the NODES
- At this stage the LMX table will be populated with the relevant routes option to each one of the nodes
- The output of the LMX will provide the Best Route
  to Reach a DLID .
  That Result Will be based on Shortest Path First (SPF)

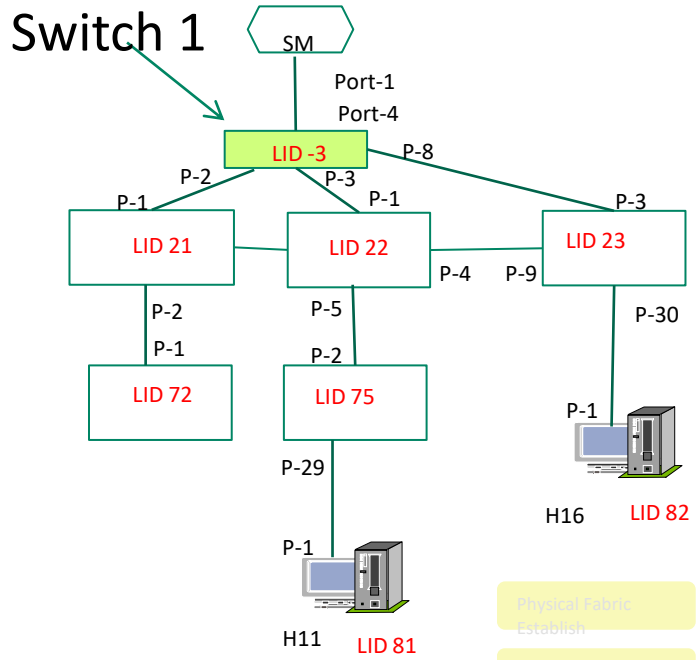

LMX Switch_1

PORT / D-LID	2	3	8	Min Hops
21	1	2	3	1
22	2	1	2	1
23	3	2	1	1
75	3	2	3	2
81	4	3	4	3
82	4	3	2	2

LFT Switch_1

The Dest. LID	Best Route / exit port
21	2
22	3
23	8
75	3
81	3
82	8

ADAPTIVE ROUTING (DYNAMIC)

- AR PORT GROUP – AR will automatically group switch multiple ports , Having the same cost (minhop) towards any specific destination lid
- Allows a switch ,to move the data connection between exit ports, selecting the least congested port of that AR port group
- The best exit port to be used , is analyzed and selected by a port transmit “virtual que manager”
- AR functionality is managed by a new SM component called Adaptive Routing Manager
  (AR plugin)

AR Group	Ports
1	{25,27}

Dest. LID	AR Group
5	1
6	1
7	1
8	1
9	1

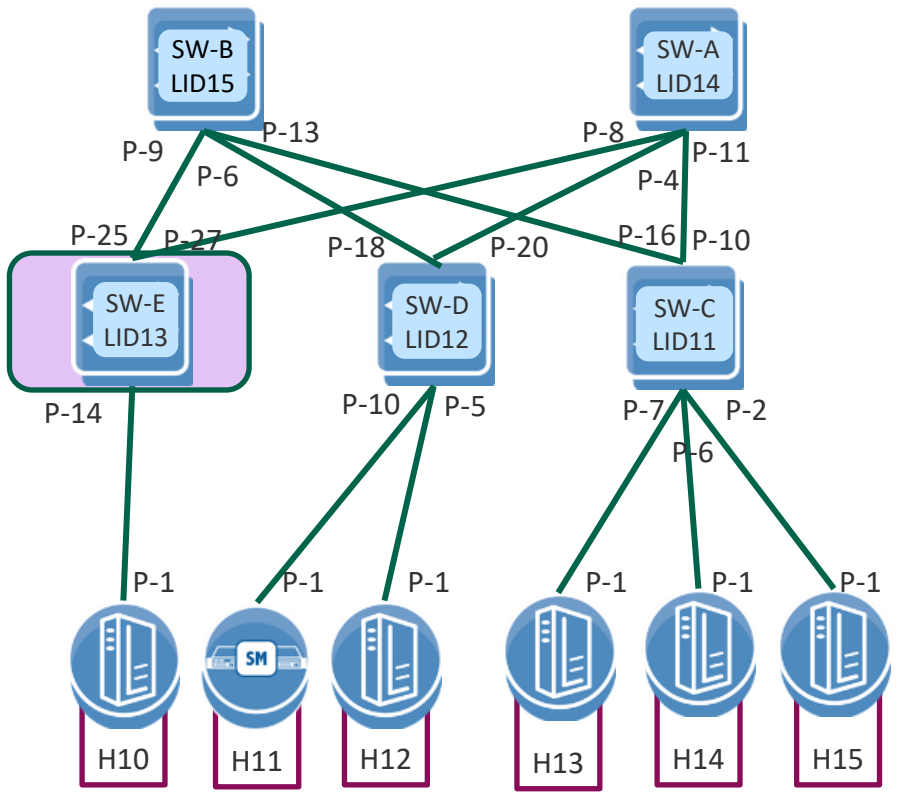

Adaptive Routing

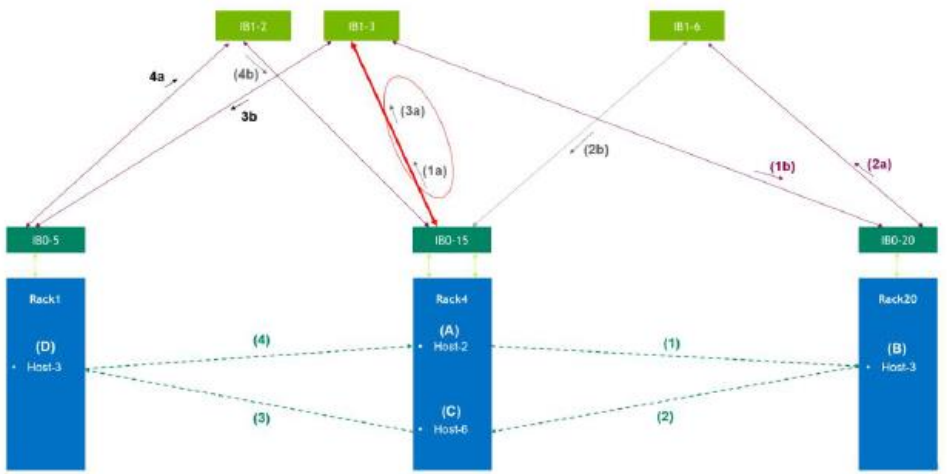

Communication paths during NCCL AllReduce

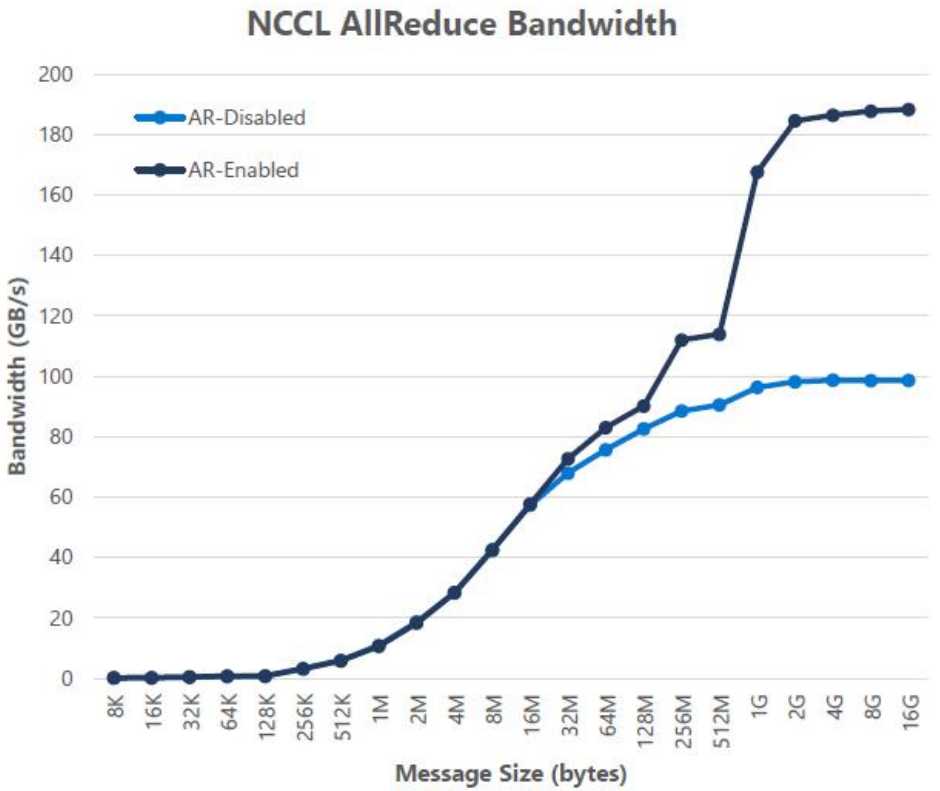

Impact of Adaptive Routing

- Congestion can happen with static routing if a single link is being used by two or more communicating pairs
- AR avoids congestion and offers stable performance

SHIELD

- Regular forwarding decision is a single output port per DLID.
  - Example Forwarding Data Base (FDB) for Sw00:

Dlid	Output Port	Port Status
Dest0	0	Up
Dest1	1	Down

  - When a link fails, traffic sent over is discarded.
- With SHIELD and FRNs, other link options are made available and used if needed.
  - Example FDB with AR for Sw00:

Dlid	Output Port	Port Status
Dest0	0,1	Up, Down
Dest1	0,1	Up, Down

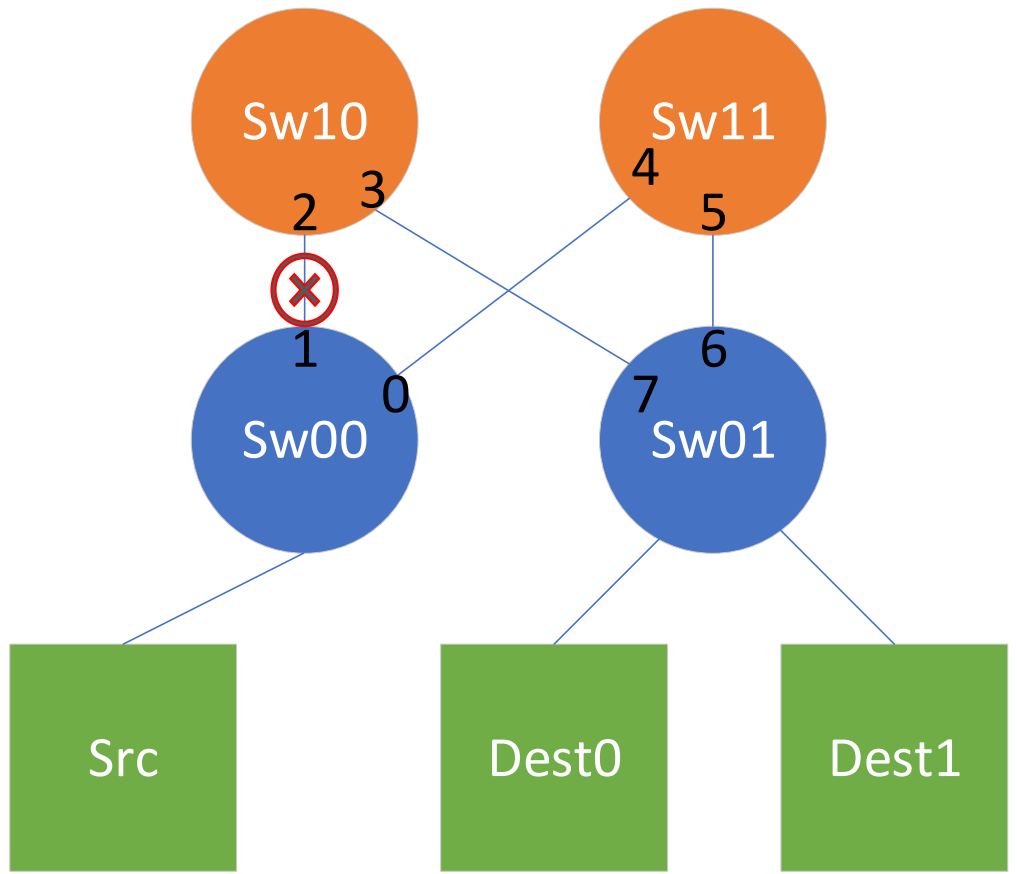

OPENSM

	Switch SM	UFM SM	OFED SM
Cost	• Free	• Not Free	• Free
HA	• Opensm HA	• UFM HA	• None HA
Failover Speed	• Medium	• Fast	• Normal
Version	• Less	• More	• Normal
Feature	• Old	• Newest	• Normal
Configuration	• Very Difficult	• Easy	• Difficult

CLOUD NETWORKING API

Allows operators to create/remove/update tenants

Manages PKey GUIDs by getting, adding, and removing GUIDs from PKeys

Isolate tenant networks

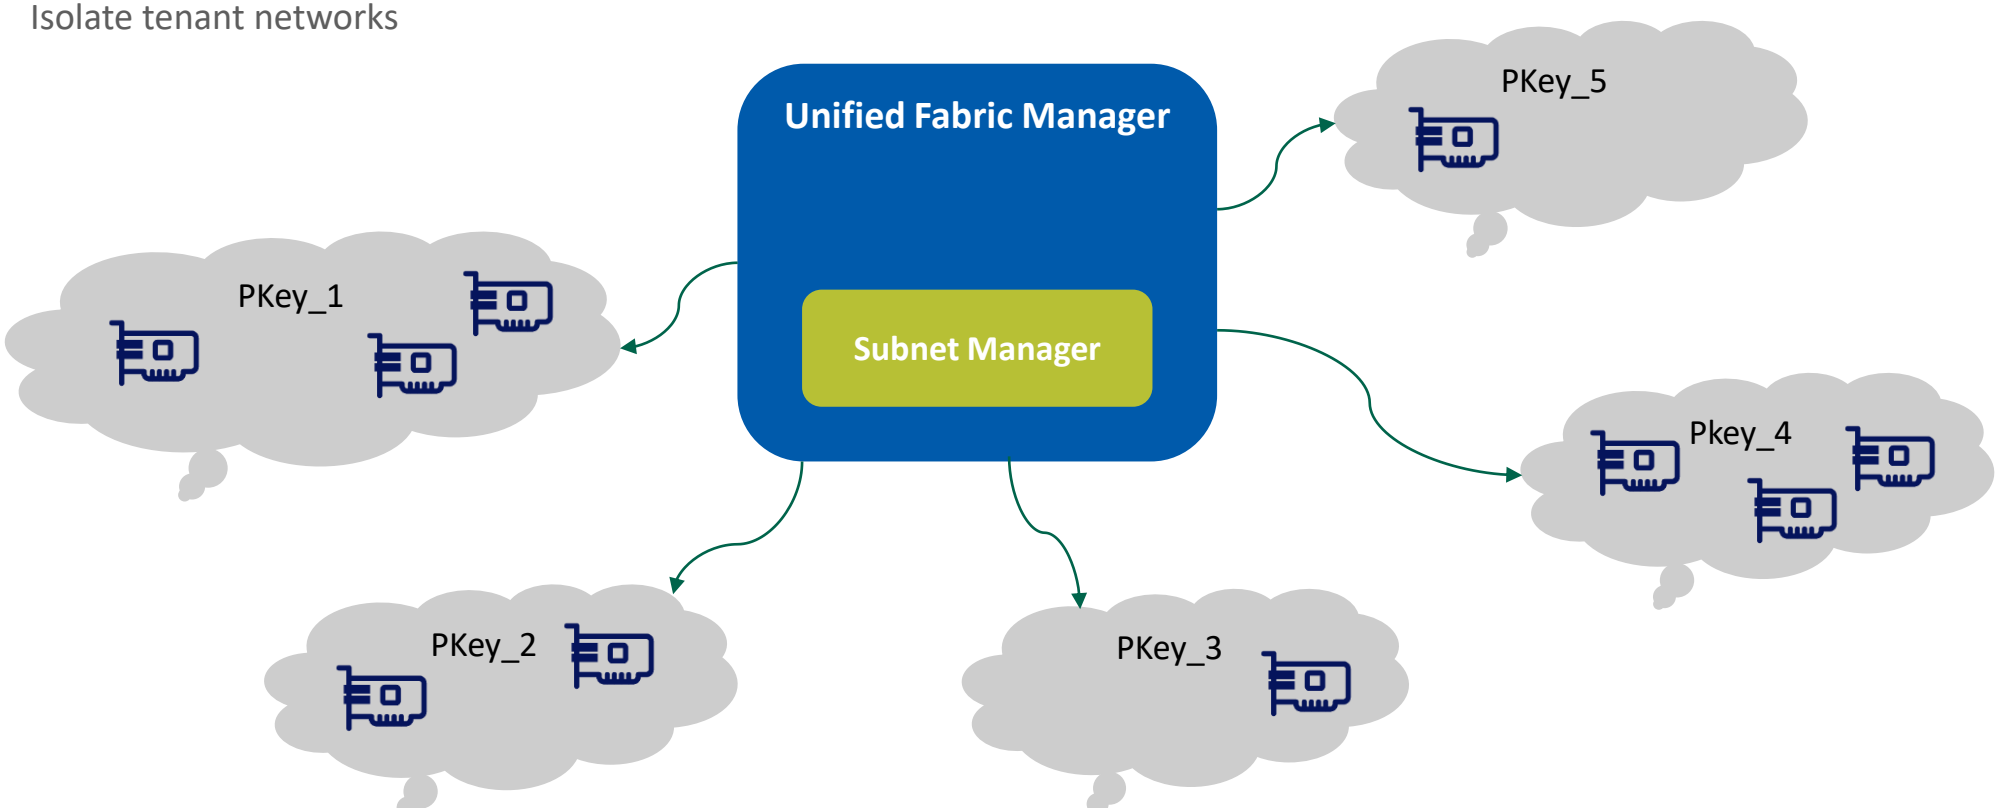

IB NETWORK AUTOMATION & PROVISIONING

- ✓ Default Network
  - ✓ FULL - Only SM and Storage
  - ✓ Limited - All nodes
- ✓ Tenant Network
  - ✓ Full - All nodes

File	Required configuration in UFM
partition.conf	Default=0x7fff, all=limited, self=FULL, IO-Nodes-GUID=FULL TenantA=0x8001, all=full TenantB=0x8002, all=full …

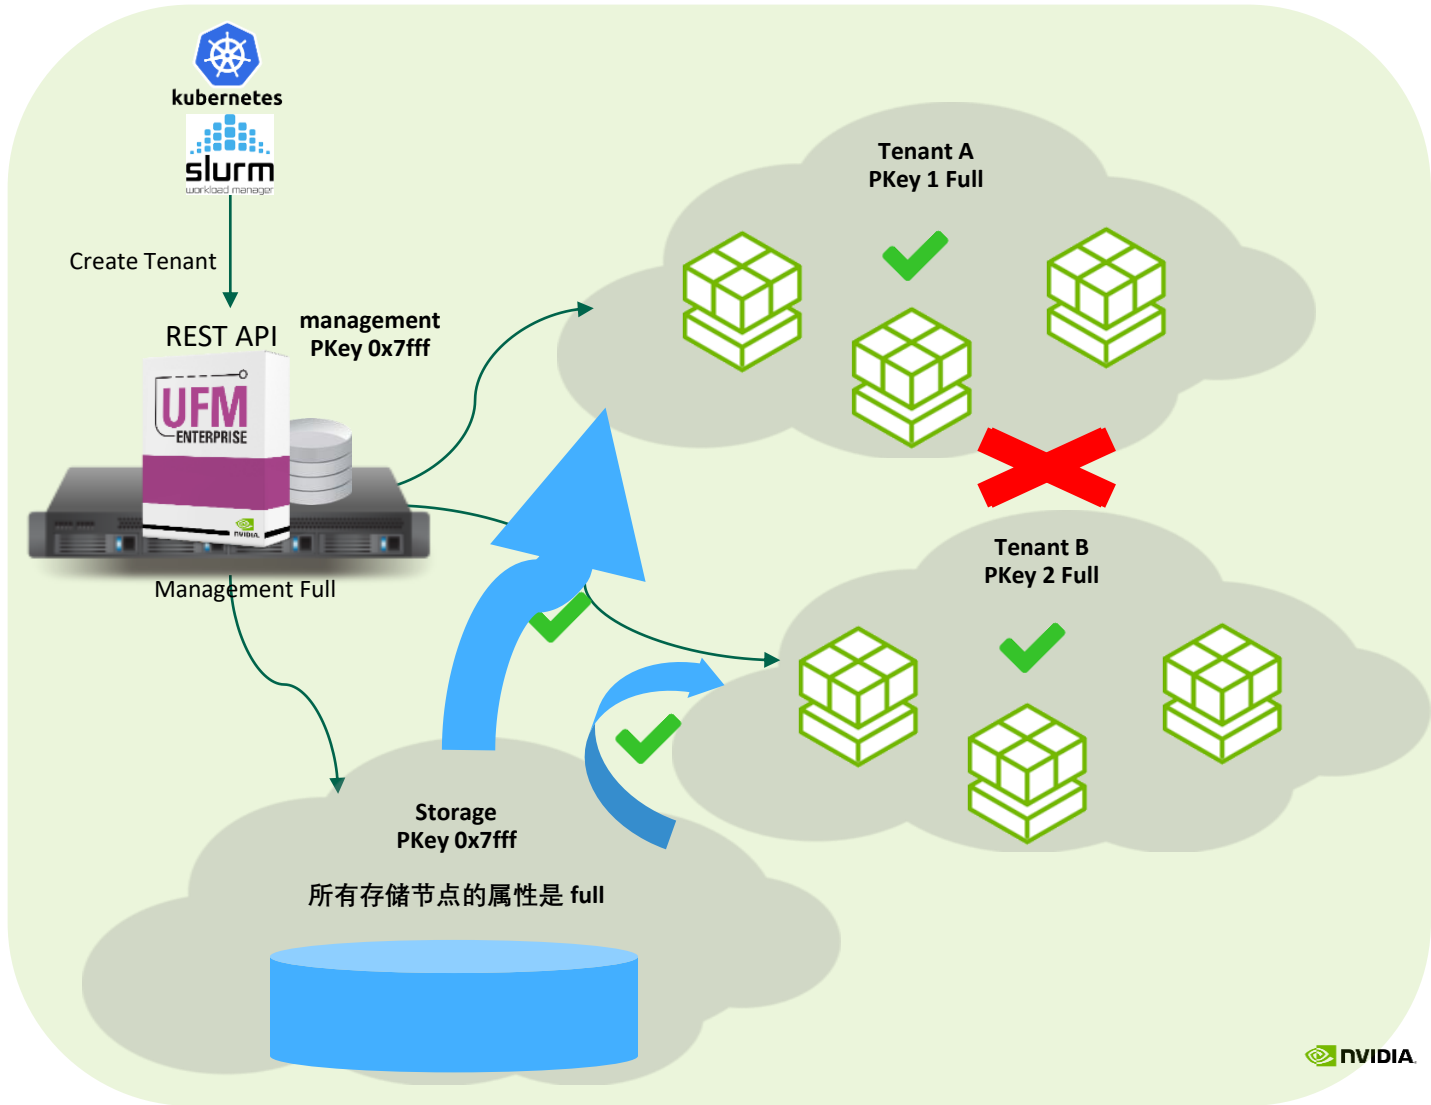

INFINIBAND OPENSTACK - ORCHESTRATING HPC CLOUDS

- Native InfiniBand integration into OpenStack
- RDMA-enabled virtual machines
- Network isolation and partition
- Cluster management with UFM appliance
- InfiniBand In-Network Computing
- Accelerate cloud storage NVMe over Fabrics

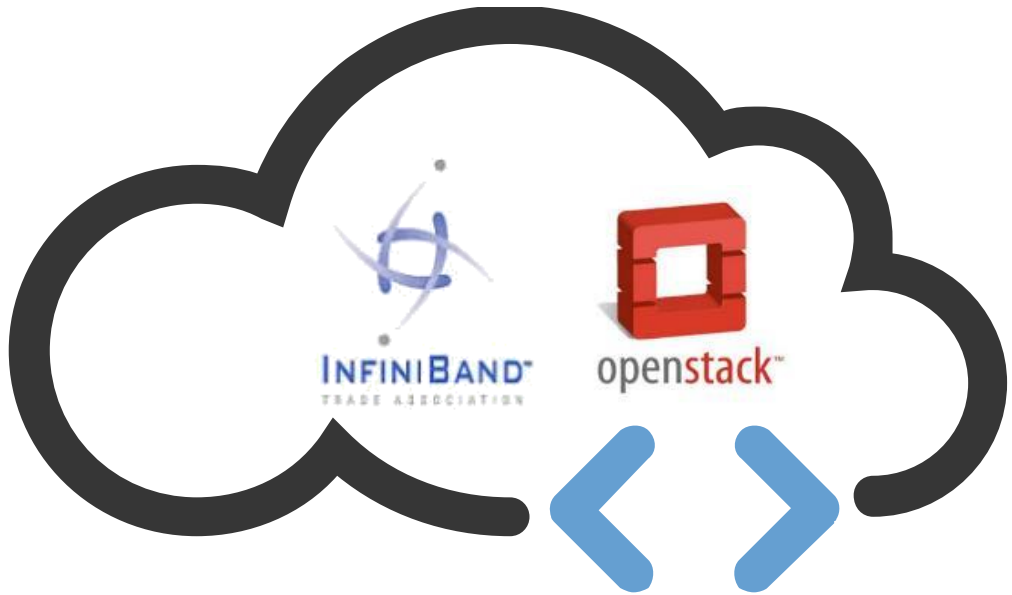

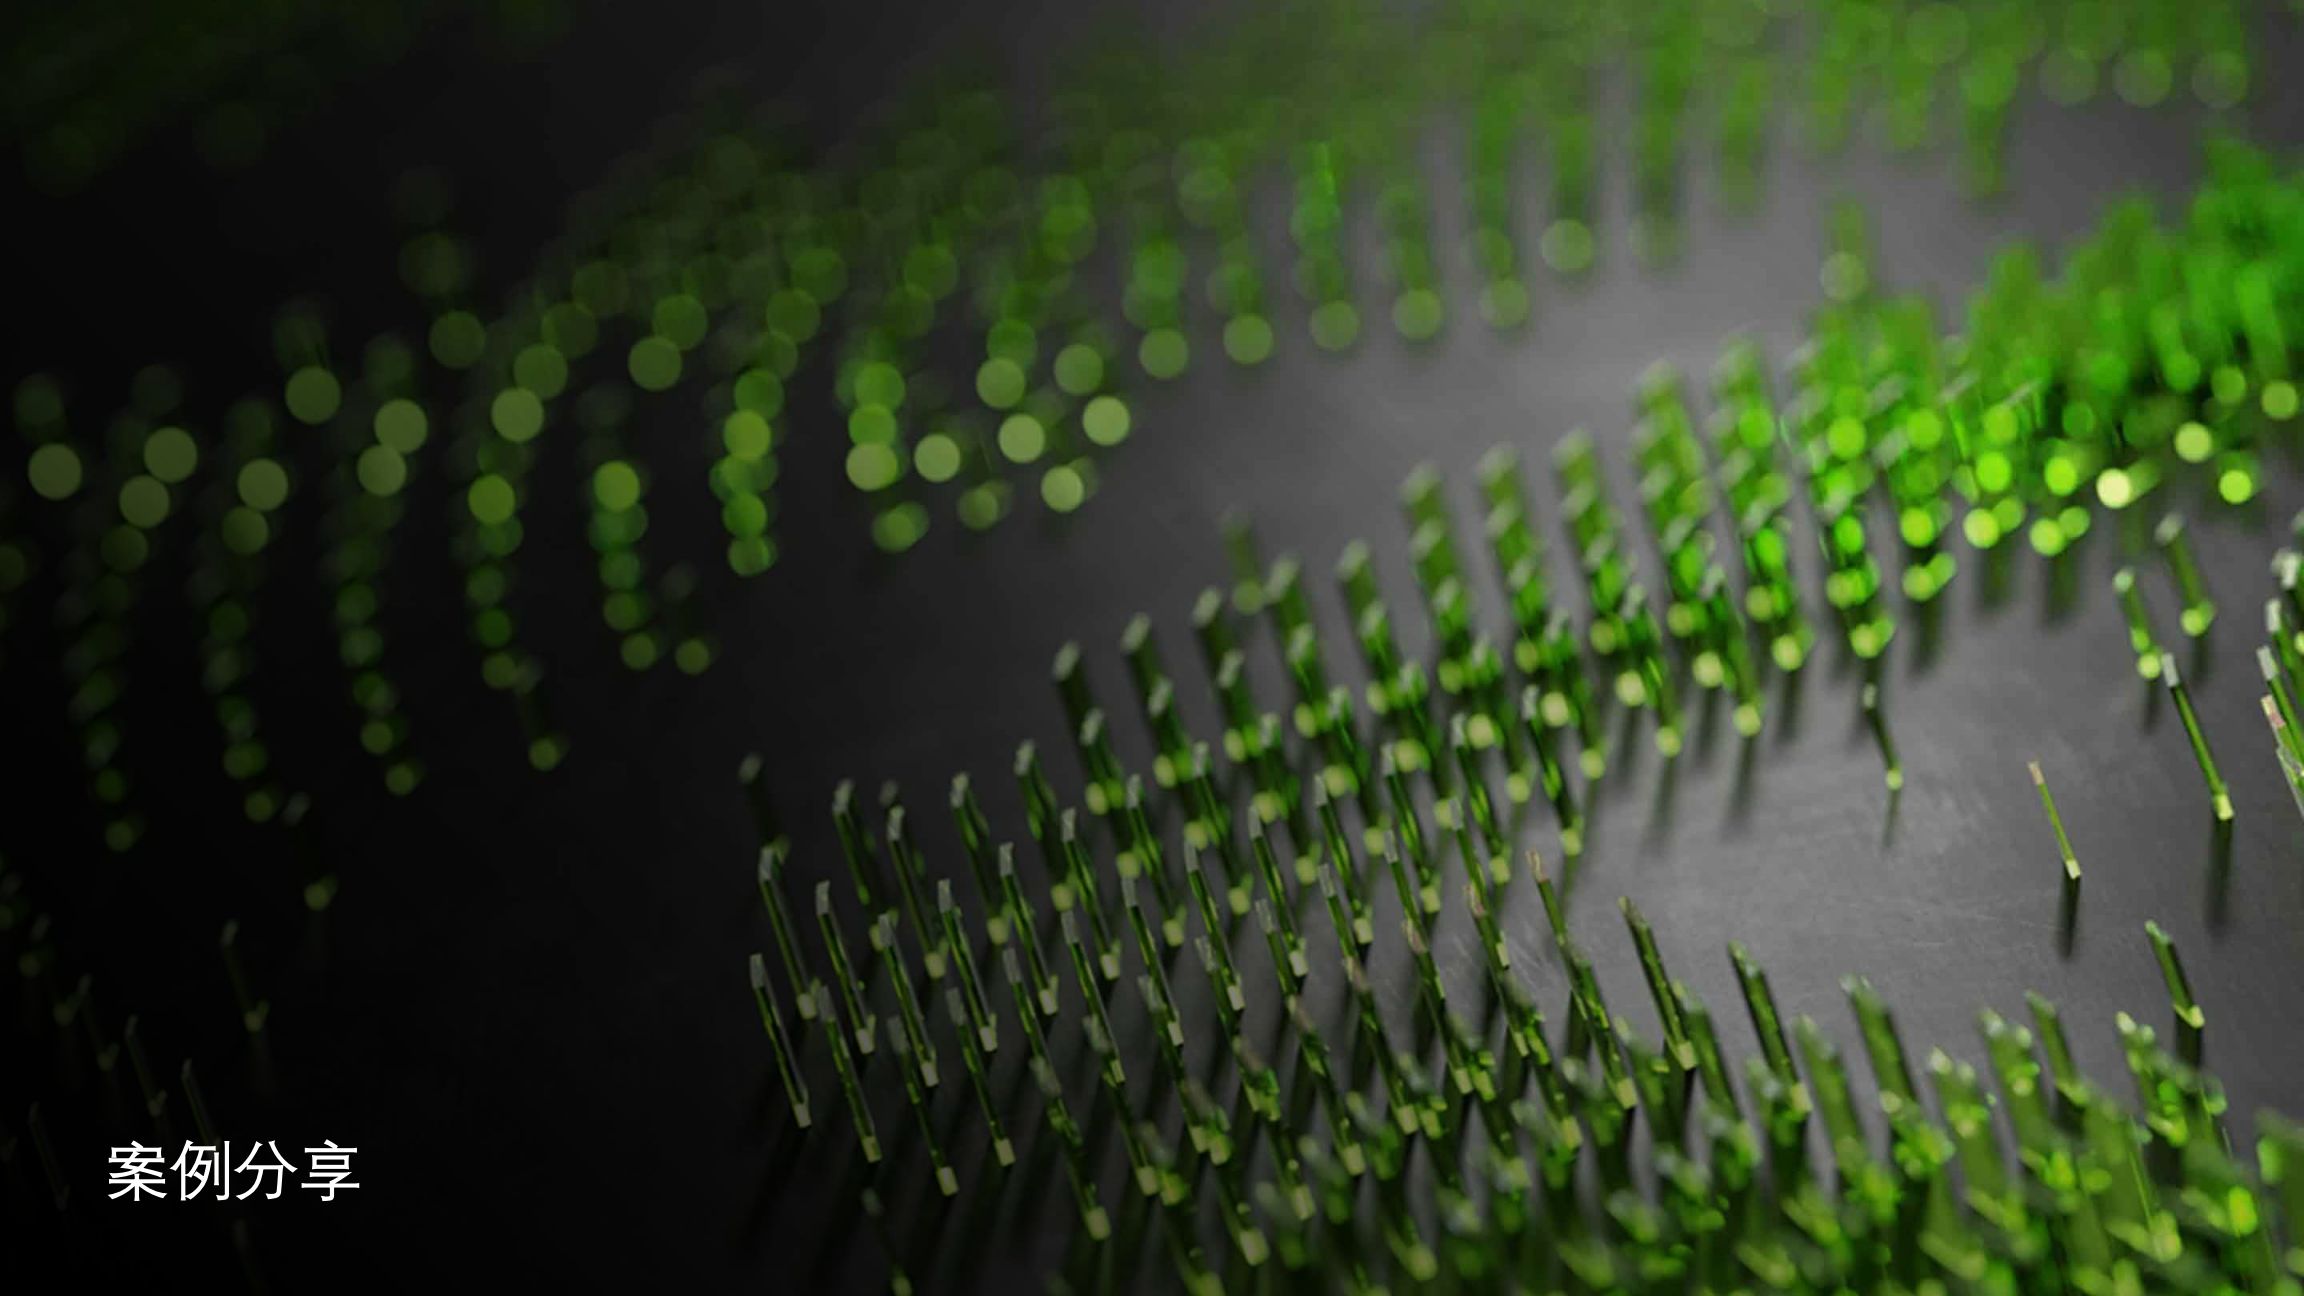
案例分享

INFINIBAND CLOUD

不安全 | doit.com.cn/p/348468.html

200G HDR InfiniBand 加速 Microsoft Azure HPC云

宋 家雨 发布于 2019-11-25　　分类：业界

微软公司 Azure 计算事业部副总裁 Girish Bablani 表示：“Microsoft Azure 旨在为寻求于云中运行计算和数据密集型应用程序的客户带来领先的性能和可扩展性。此外,我们还努力确保客户可使用在其本地超级计算机上运行的相同软件驱动程序和库。借助 200 GB HDR InfiniBand,我们能够为真实的 HPC 和人工智能工 metal 超级计算机相媲美的可扩展性与性能。”

Microsoft Azure

Mellanox Quantum

SHARP 2

OpenAI

server.zhiding.cn/server/2020/0521/3126489.shtml

本周微软宣布，已经在Azure云中托管了OpenAI排名第五的AI超级计算机。2019年微软向OpenAI行业研究小组投资了10亿美元。

本周微软宣布，已经在Azure云中托管了OpenAI排名第五的AI超级计算机。2019年微软向OpenAI行业研究小组投资了10亿美元。

本周微软宣布，已经在Azure云中托管了OpenAI排名第五的AI超级计算机。2019年微软向OpenAI行业研究小组投资了10亿美元。这个AI超算系统包括大约10000个GPU和285000多个CPU核心，将用于提升处理超大型AI模型的能力，据OpenAI称，大型AI模型的规模每3.5个月就会翻一番。微软用于自然语言生成的Turing模型包含约170亿个参数，比去年的最大模型增加了17倍。因此，这个超级计算机将大有用处。

Achieve more with Azure HPC

Purpose-built HPC	A full range of CPU and GPU capabilities that help applications scale to 80K+ cores
Fast, secure networking	Fast InfiniBand inter-connects as well as edge-to-cloud connectivity
High performing storage	A range of storage capabilities to support simple-to-complex storage needs
Workload orchestration	End-to-end workflow agility using known, familiar tools and processes
Intelligence services	AI, machine learning, and deep learning at supercomputer scale

ANSYS FLUENT ON HBV2

- **App:** ANSYS Fluent
- **Version:** 14.06.004
- **Model:** External Flow over a Formula-1 Race Car (f1_racecar_140m)
- **Configuration Details:** 60 MPI ranks were run (2 out of 4 cores per NUMA) in each HBv2 VM in order to leave nominal resources to run Linux background processes and give ~6 GB/s of memory bandwidth per core. In addition, Adaptive Routing was enabled and DCT (Dynamic Connected Transport) was used as the transport layer, while **HPC-X version 2.50 (UCX v1.6) was used for MPI**. Azure CentOS HPC 7.6 image was used from https://github.com/Azure/azhpc-images
- **Summary:** HBv2 VMs scale super linearly (112%) up to the top end measured number of VMs (128). The Fluent Solver Rating measured at this top-end level of scale is 83% more performance than the current leader submission on ANSYS public database for this model (https://bit.ly/2OdAExM).

Microsoft Azure

STAR-CCM + ON HBV2

App: Siemens Star-CCM+
Version: 14.06.004
Model: LeMans 100M Coupled Solver
Configuration Details: 116 MPI ranks were run (4 ranks from each of 29 NUMA) in each HBv2 VM in order to leave nominal resources to run Linux background processes. In addition, Adaptive Routing was enabled and DCT (Dynamic Connected Transport) was used as the transport layer, while HPC-X version 2.50 (UCX v1.6) was used for MPI. Azure CentOS HPC 7.6 image was used from https://github.com/Azure/azhpc-images

Microsoft Azure

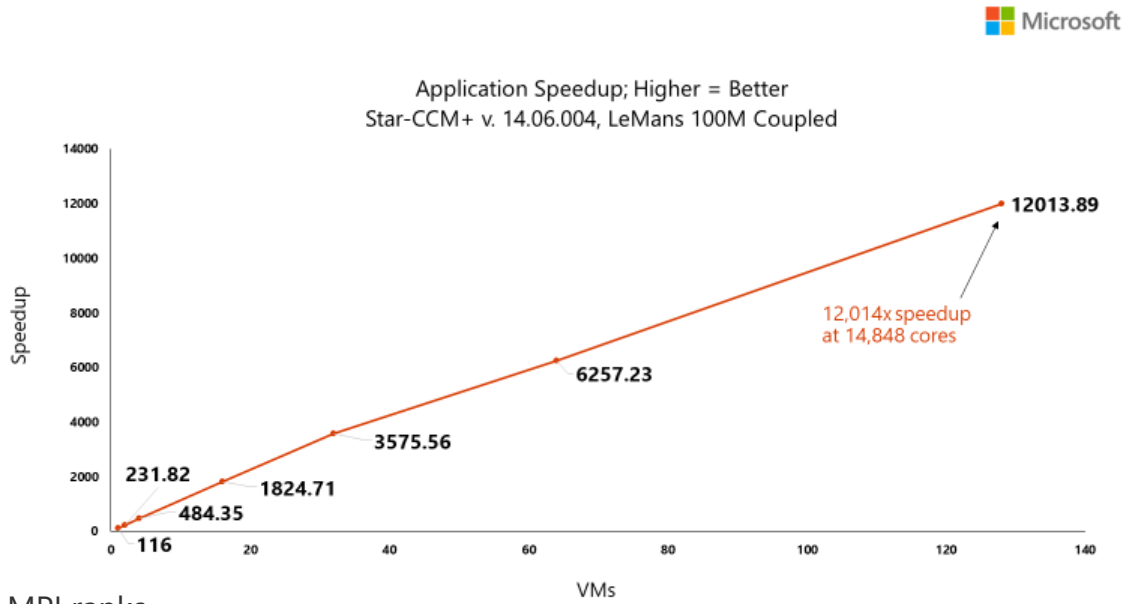

Summary: Star-CCM+ was scaled at 81% efficiency to nearly 15,000 MPI ranks delivering an application speedup of more than 12,000x. This compares favorably to Azure's previous best of more than 11,500 MPI ranks, which itself was a world-record for MPI scalability on the public cloud.

NVIDIA

VIRTUALIZATION FOR HPC: INFINIBAND SRIOV

Single Root I/O Virtualization in HPC

- **Problem**: Virtualization generally has resulted in significant I/O performance degradation (e.g., excessive DMA interrupts)
- **Solution**: SR-IOV and Mellanox InfiniBand host channel adapters
  - One physical function → multiple virtual functions, each light weight but with its own DMA streams, memory space, interrupts
  - Allows DMA to bypass hypervisor to VMs
- ***SRIOV enables virtual HPC cluster w/ near-native InfiniBand latency/bandwidth and minimal overhead***

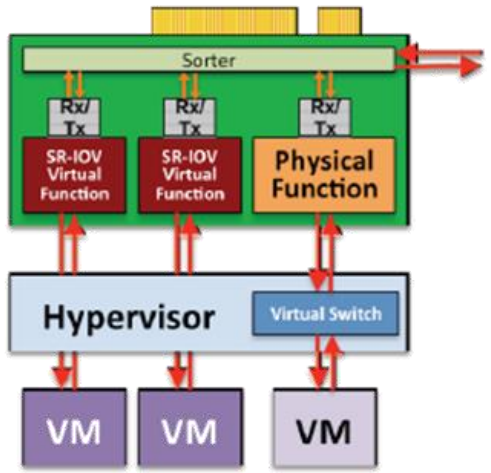

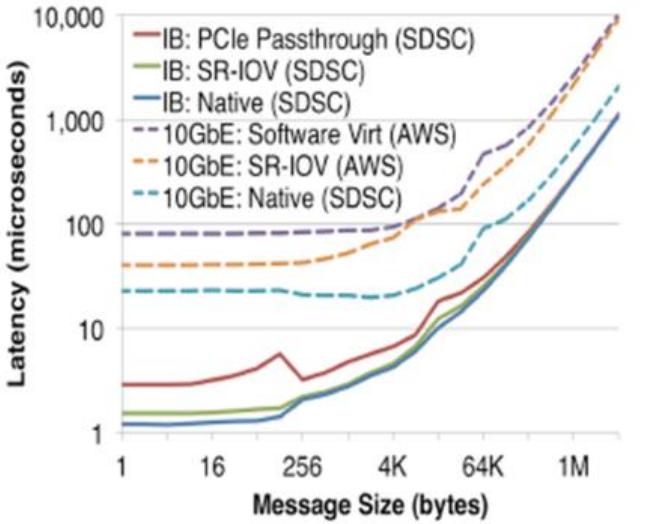

MPI point-to-point latency measured by osu_latency for QDR InfiniBand. Included for scale are the analogous 10GbE measurements from Amazon (AWS) and non-virtualized 10GbE.

SDSC SAN DIEGO SUPERCOMPUTER CENTER
at the UNIVERSITY OF CALIFORNIA; SAN DIEGO

NVIDIA

UNIVERSITY OF CAMBRIDGE: HPC CLOUD CONVERGENCE

- Motivation behind OpenStack Research Cloud
  - Make computing, data, applications and workflows more accessible, flexible and secure
  - Allow a wide range of services to be delivered from a single framework
  - Make research computing easier to use, easier to share, decrease the time to science and increasing innovation and research

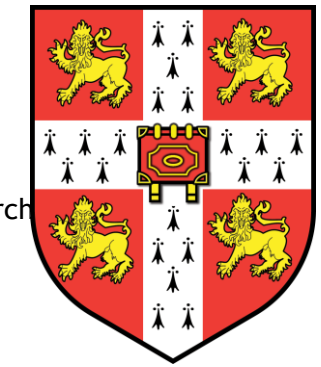

- Use cases
  - Research computing as a Service
  - HPC as a Service
  - HPDA as a Service

- Research Cloud Network Requirement
  - SRIOV and RDMA essential for HPC
  - Network virtualization with no compromise
  - High-performance data I/O

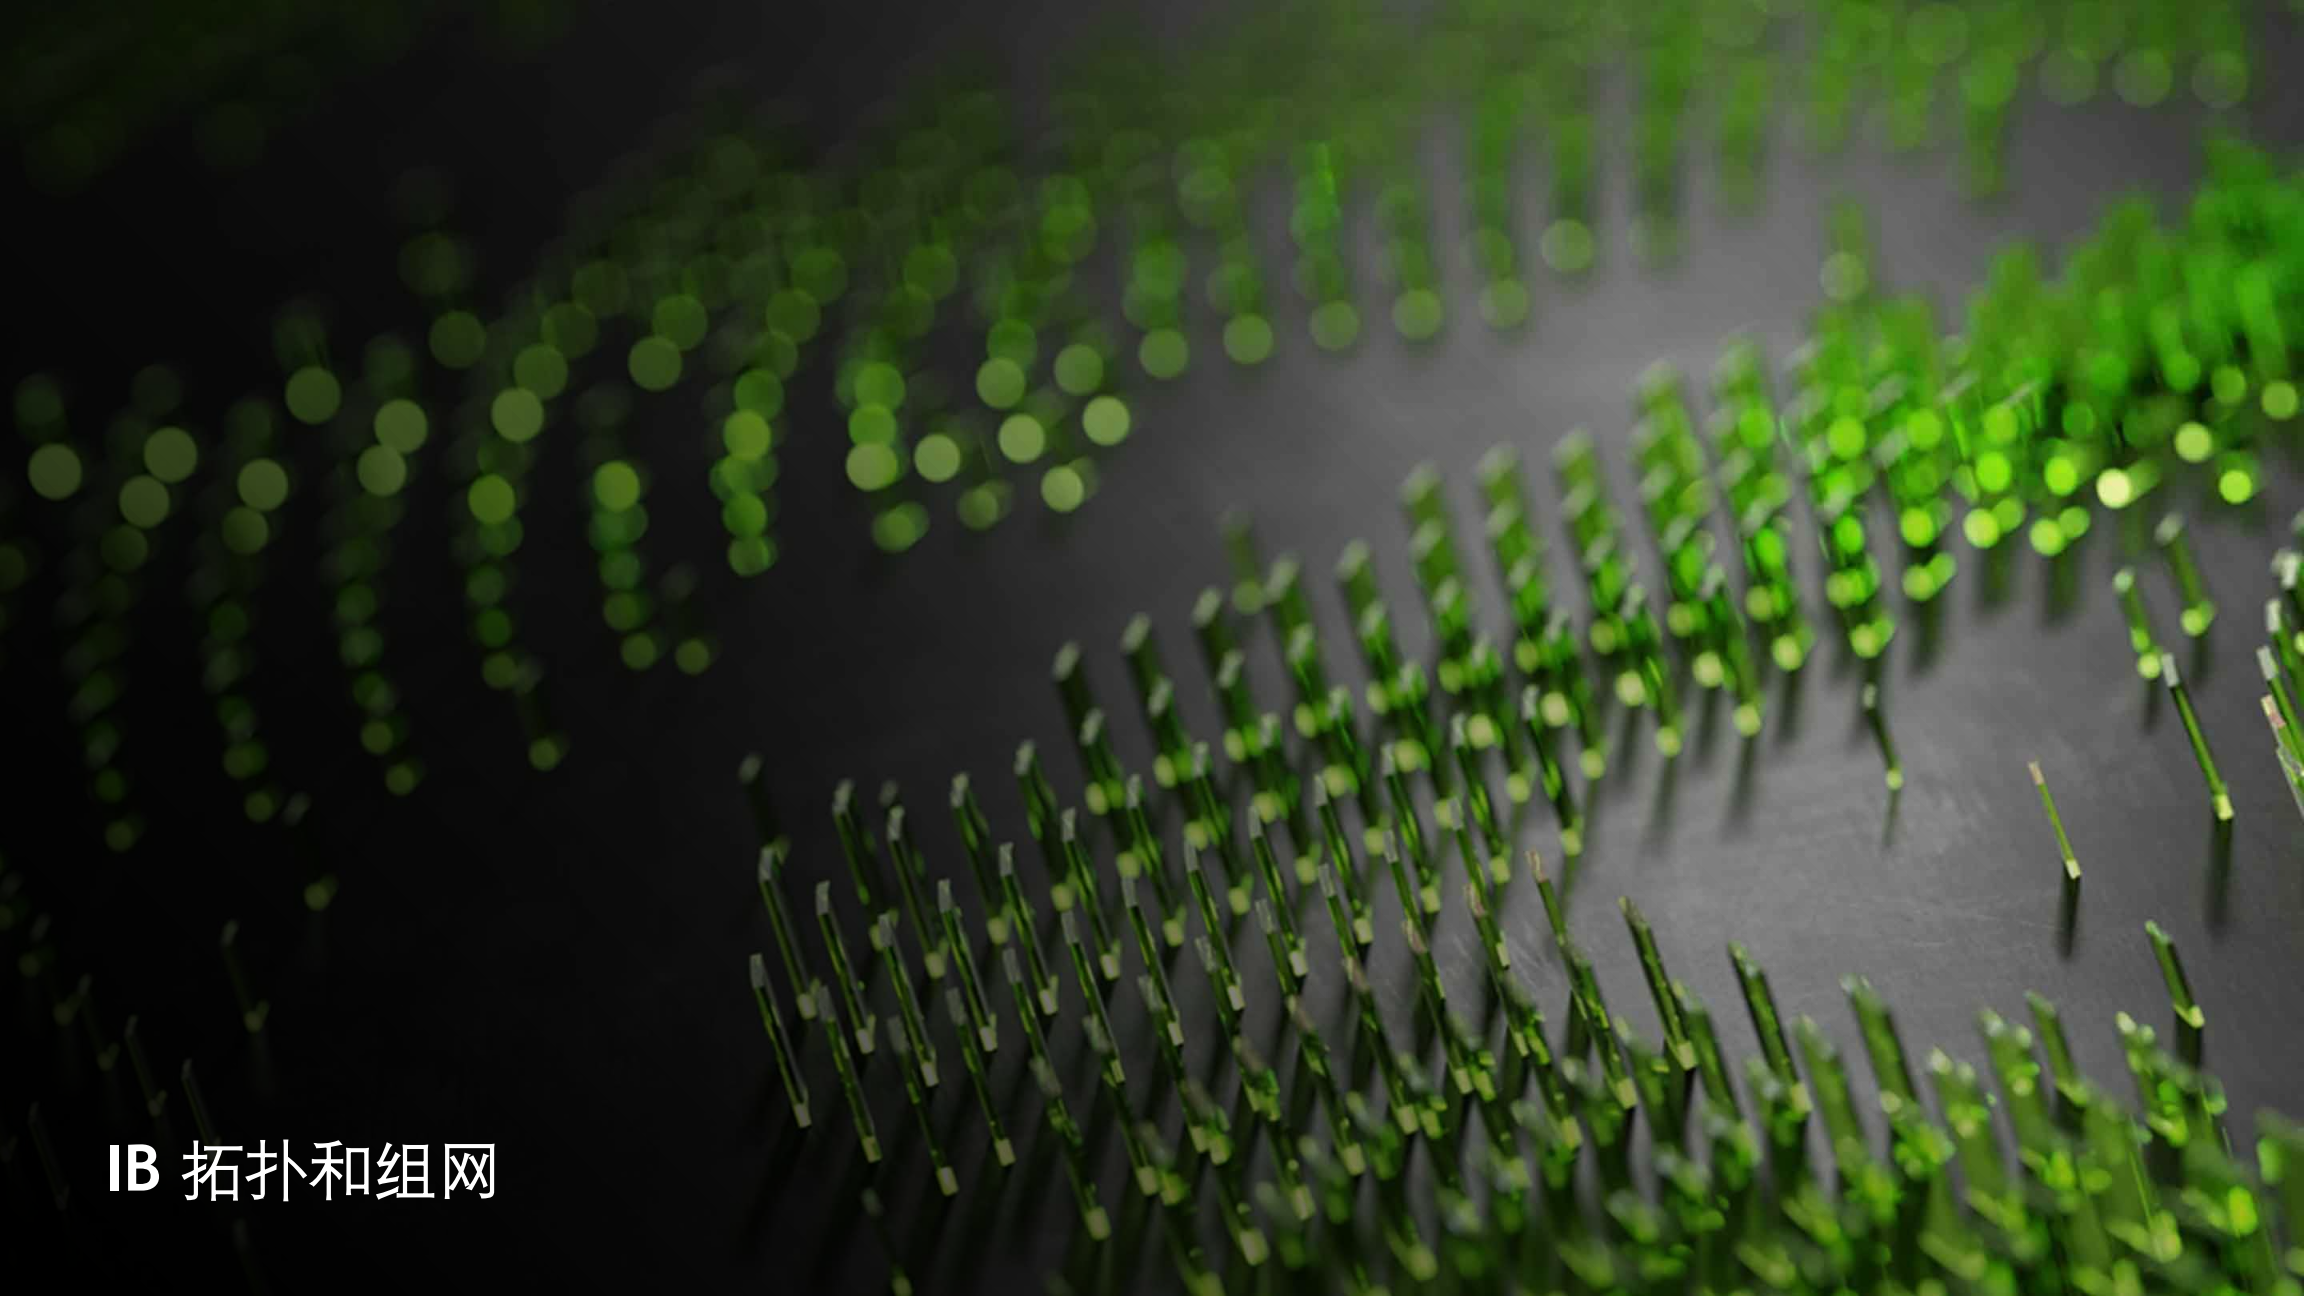

IB 拓扑和组网

SUPPORTING VARIETY OF TOPOLOGIES

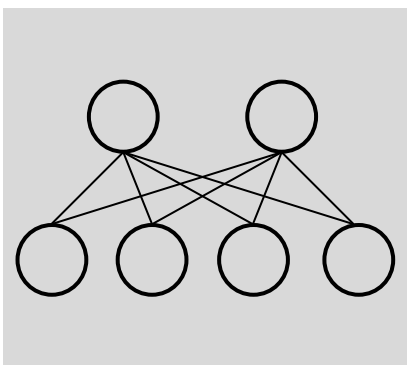

Fat Tree

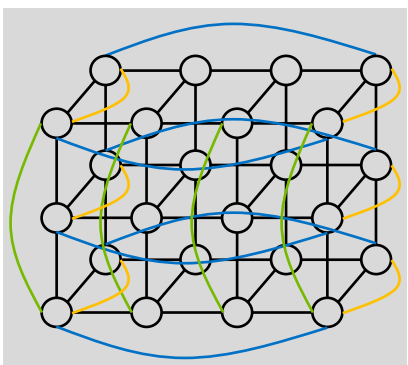

Torus

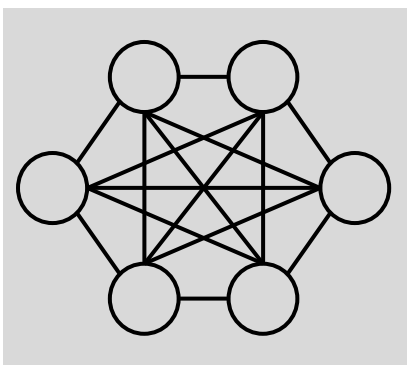

Dragonfly

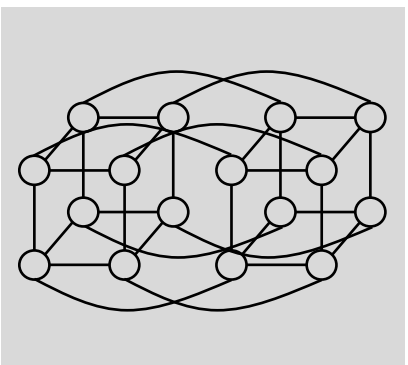

Hypercube

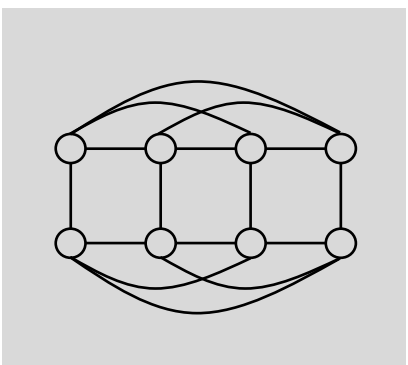

HyperX

200-Node 200G InfiniBand Platform

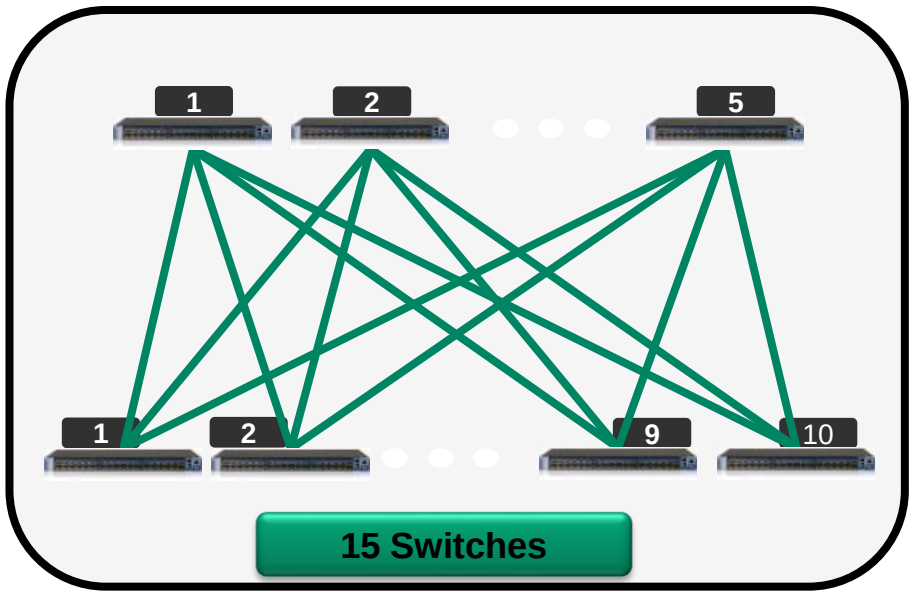

800-Node 200G InfiniBand Platform

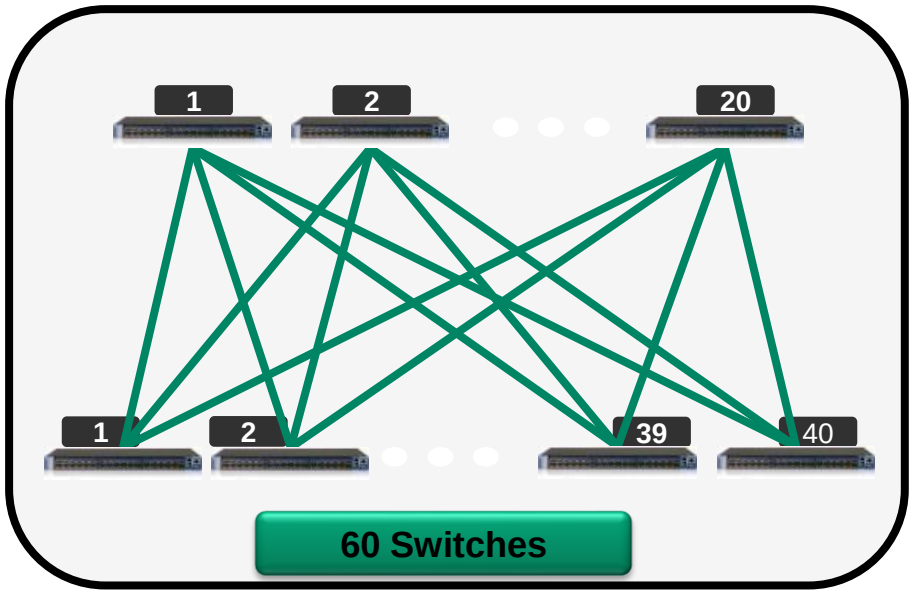

Quantum
The smartest switch, became smarter

ConnectX-6

SUPERPOD SCALABLE UNIT(SU)

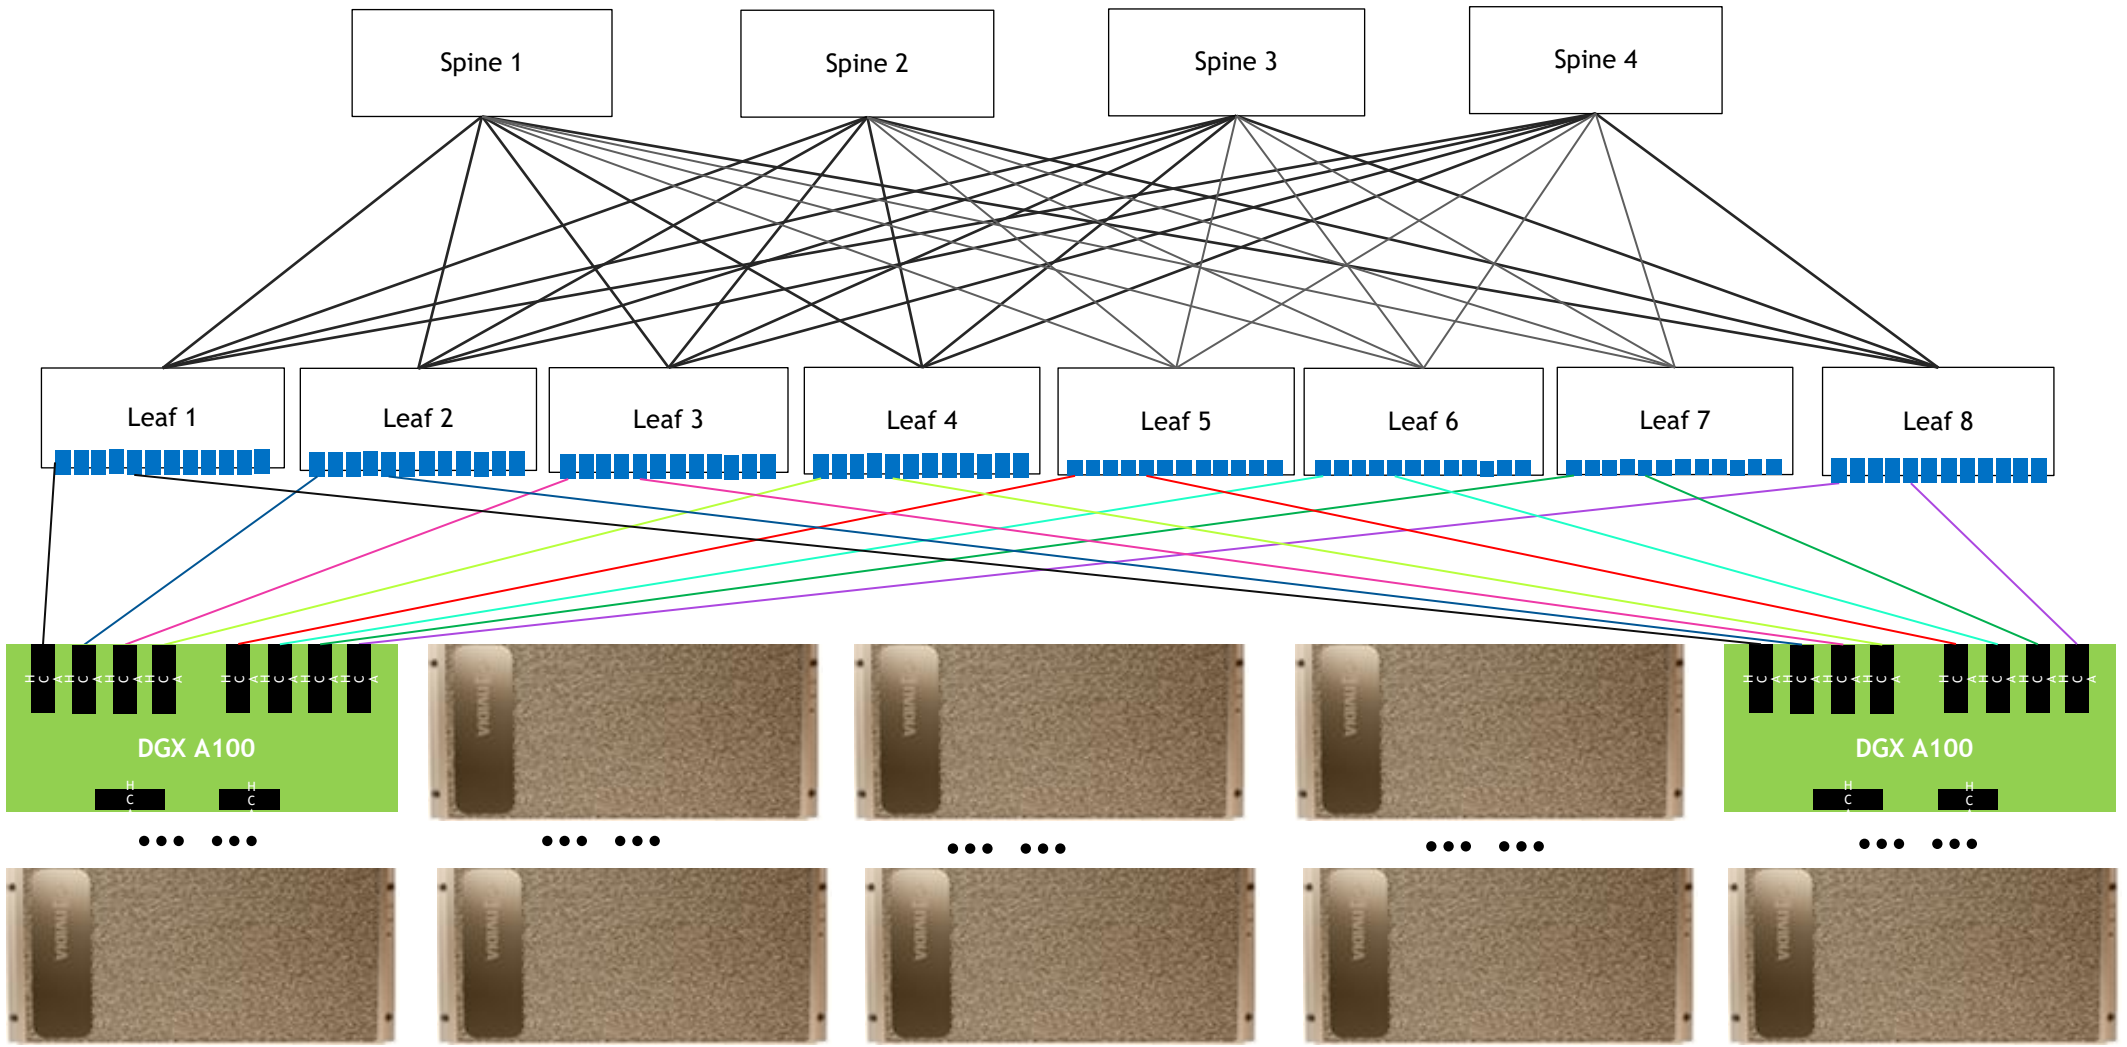

20 nodes per Scalable Unit

Example 3D Torus – SDSC Gordon

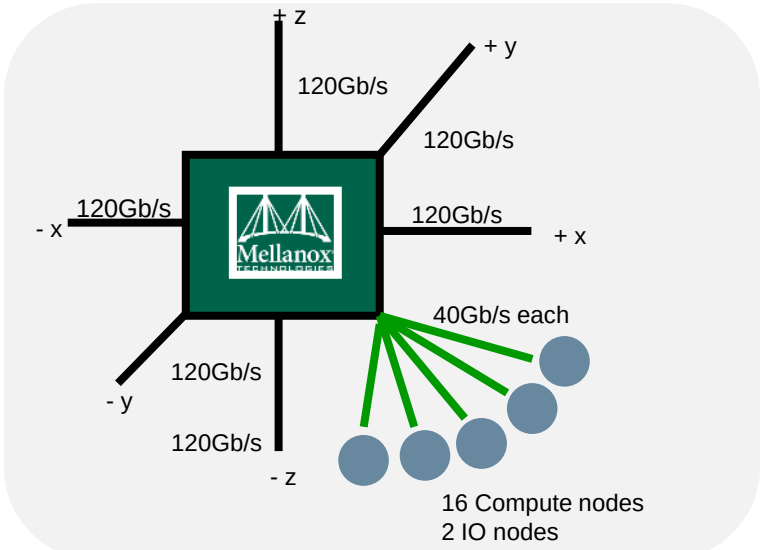

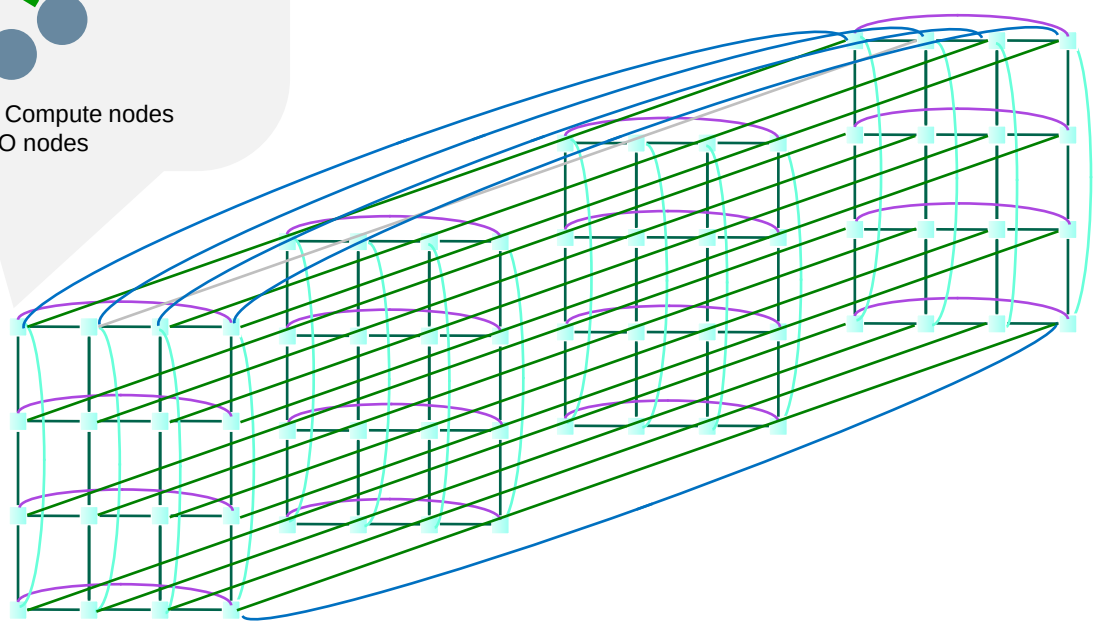

TRADITIONAL DRAGONFLY VS DRAGONFLY+

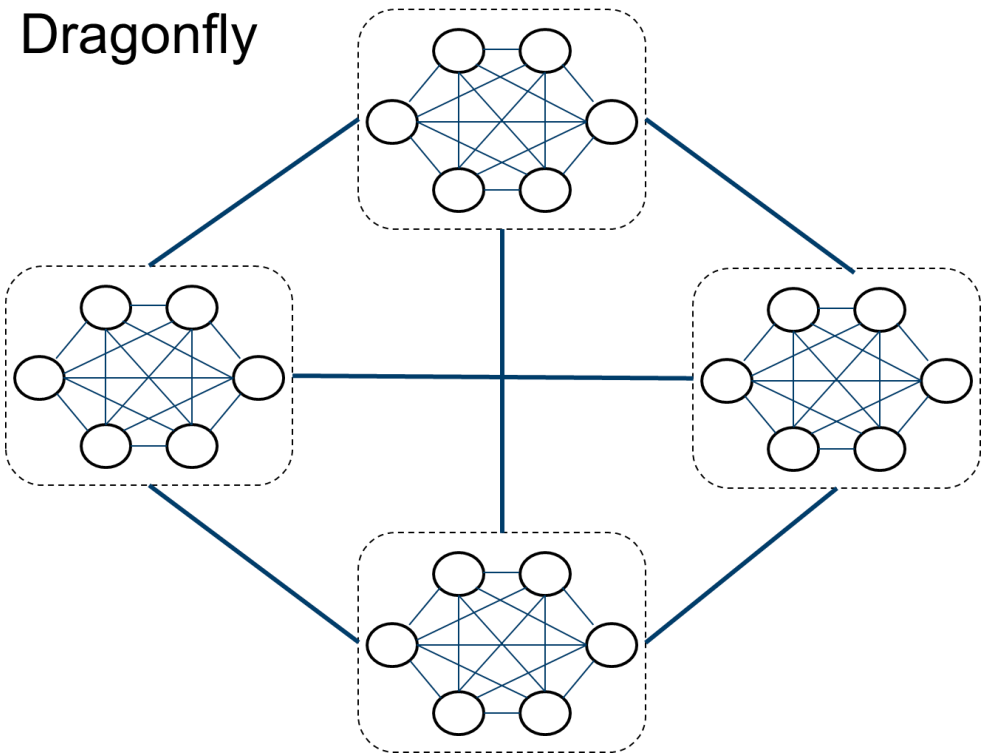

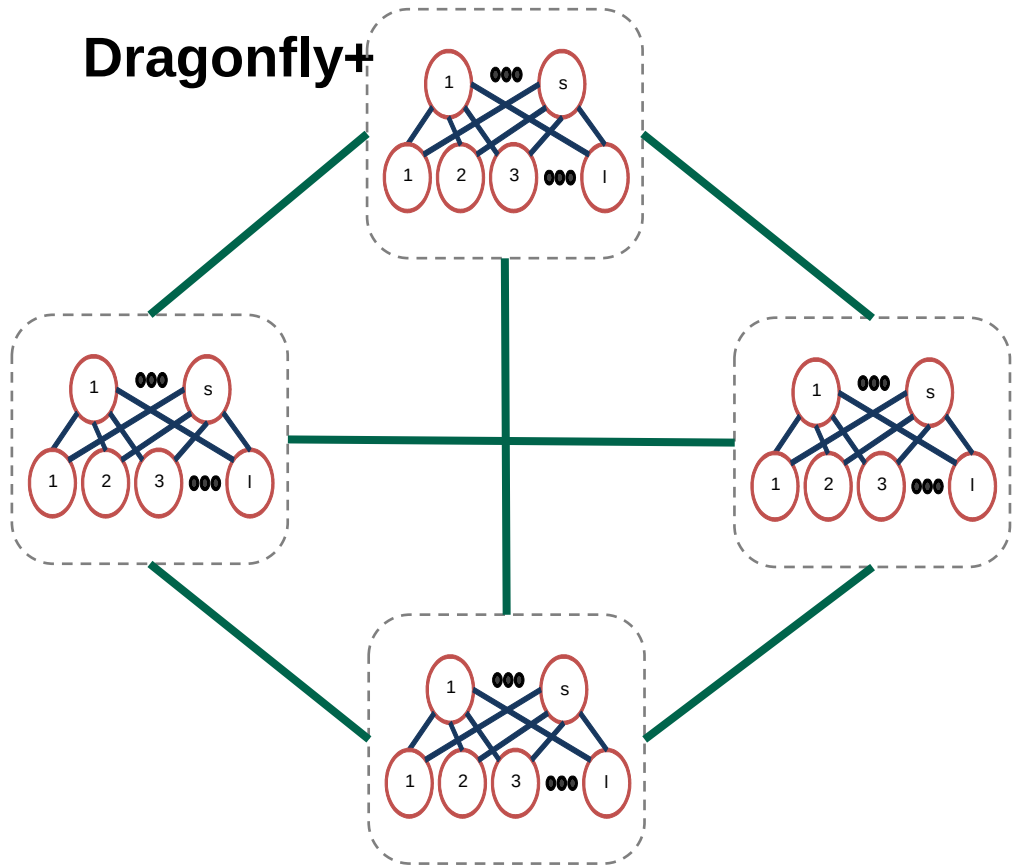

DRAGONFLY+ TOPOLOGY

Several “groups”, connected using all to all links

The topology inside each group can be any topology

Reduce total cost of network (fewer long cables)

Utilizes Adaptive Routing to for efficient operations

Simplifies future system expansion

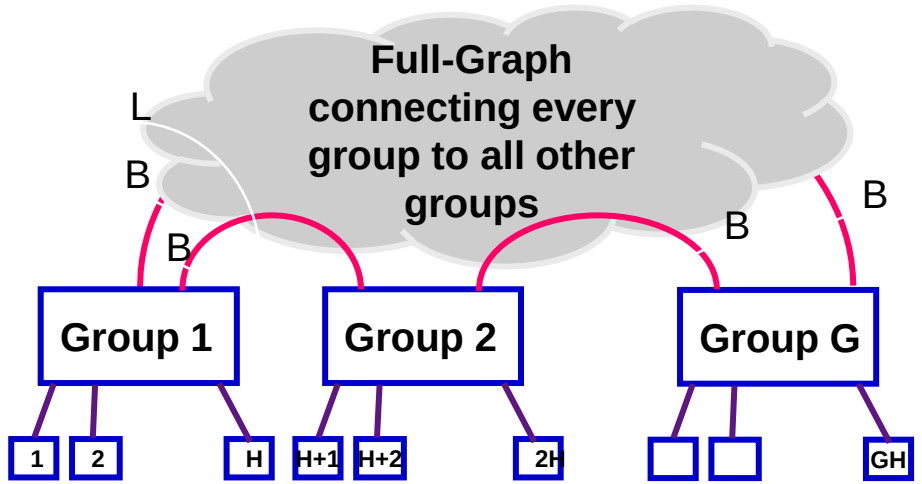

1200-Nodes Dragonfly+ Systems Example

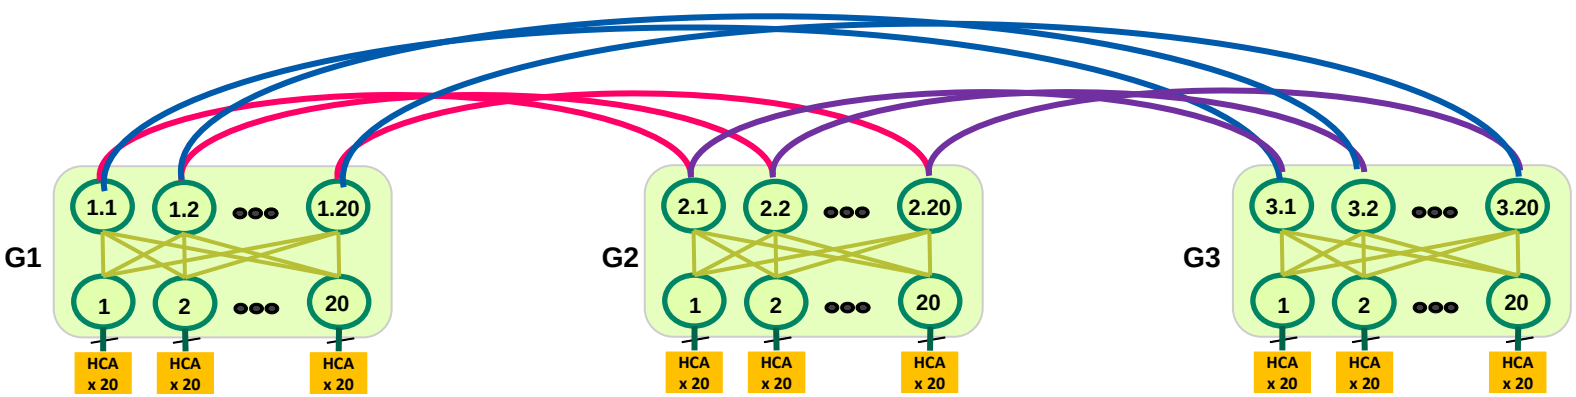

FUTURE EXPANSION OF DRAGONFLY+ BASED SYSTEM

Phase 2:
+10x400 =
8400 hosts

Re-cable the central racks,
a change local to the RACK

Phase 1:
11x400 =
4400 hosts

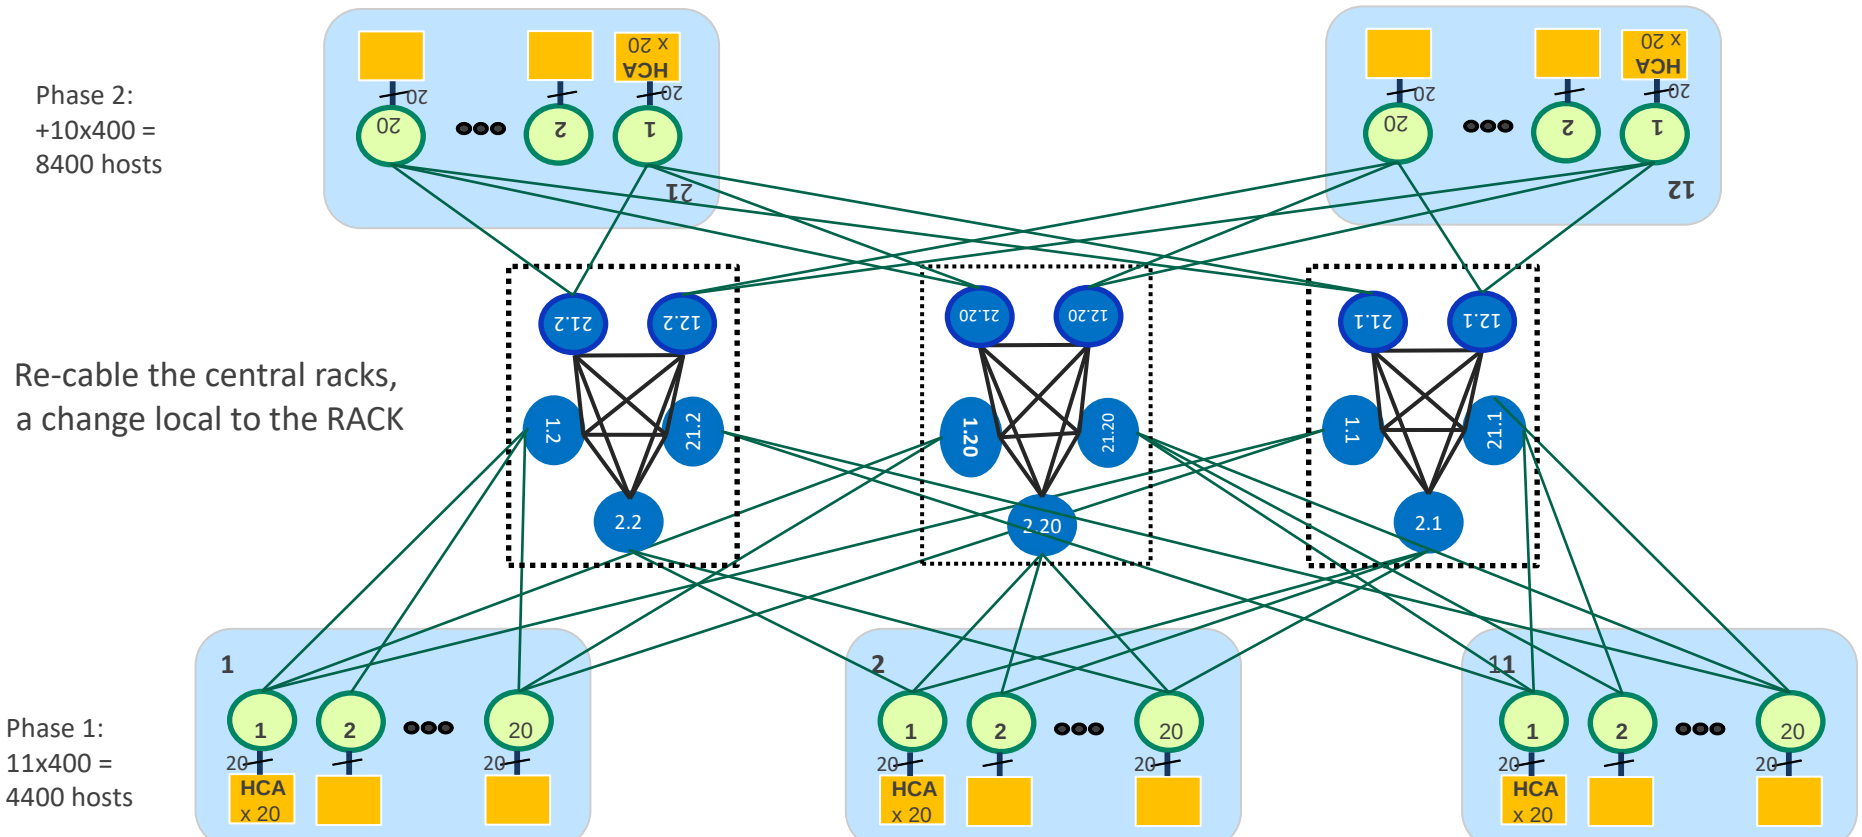

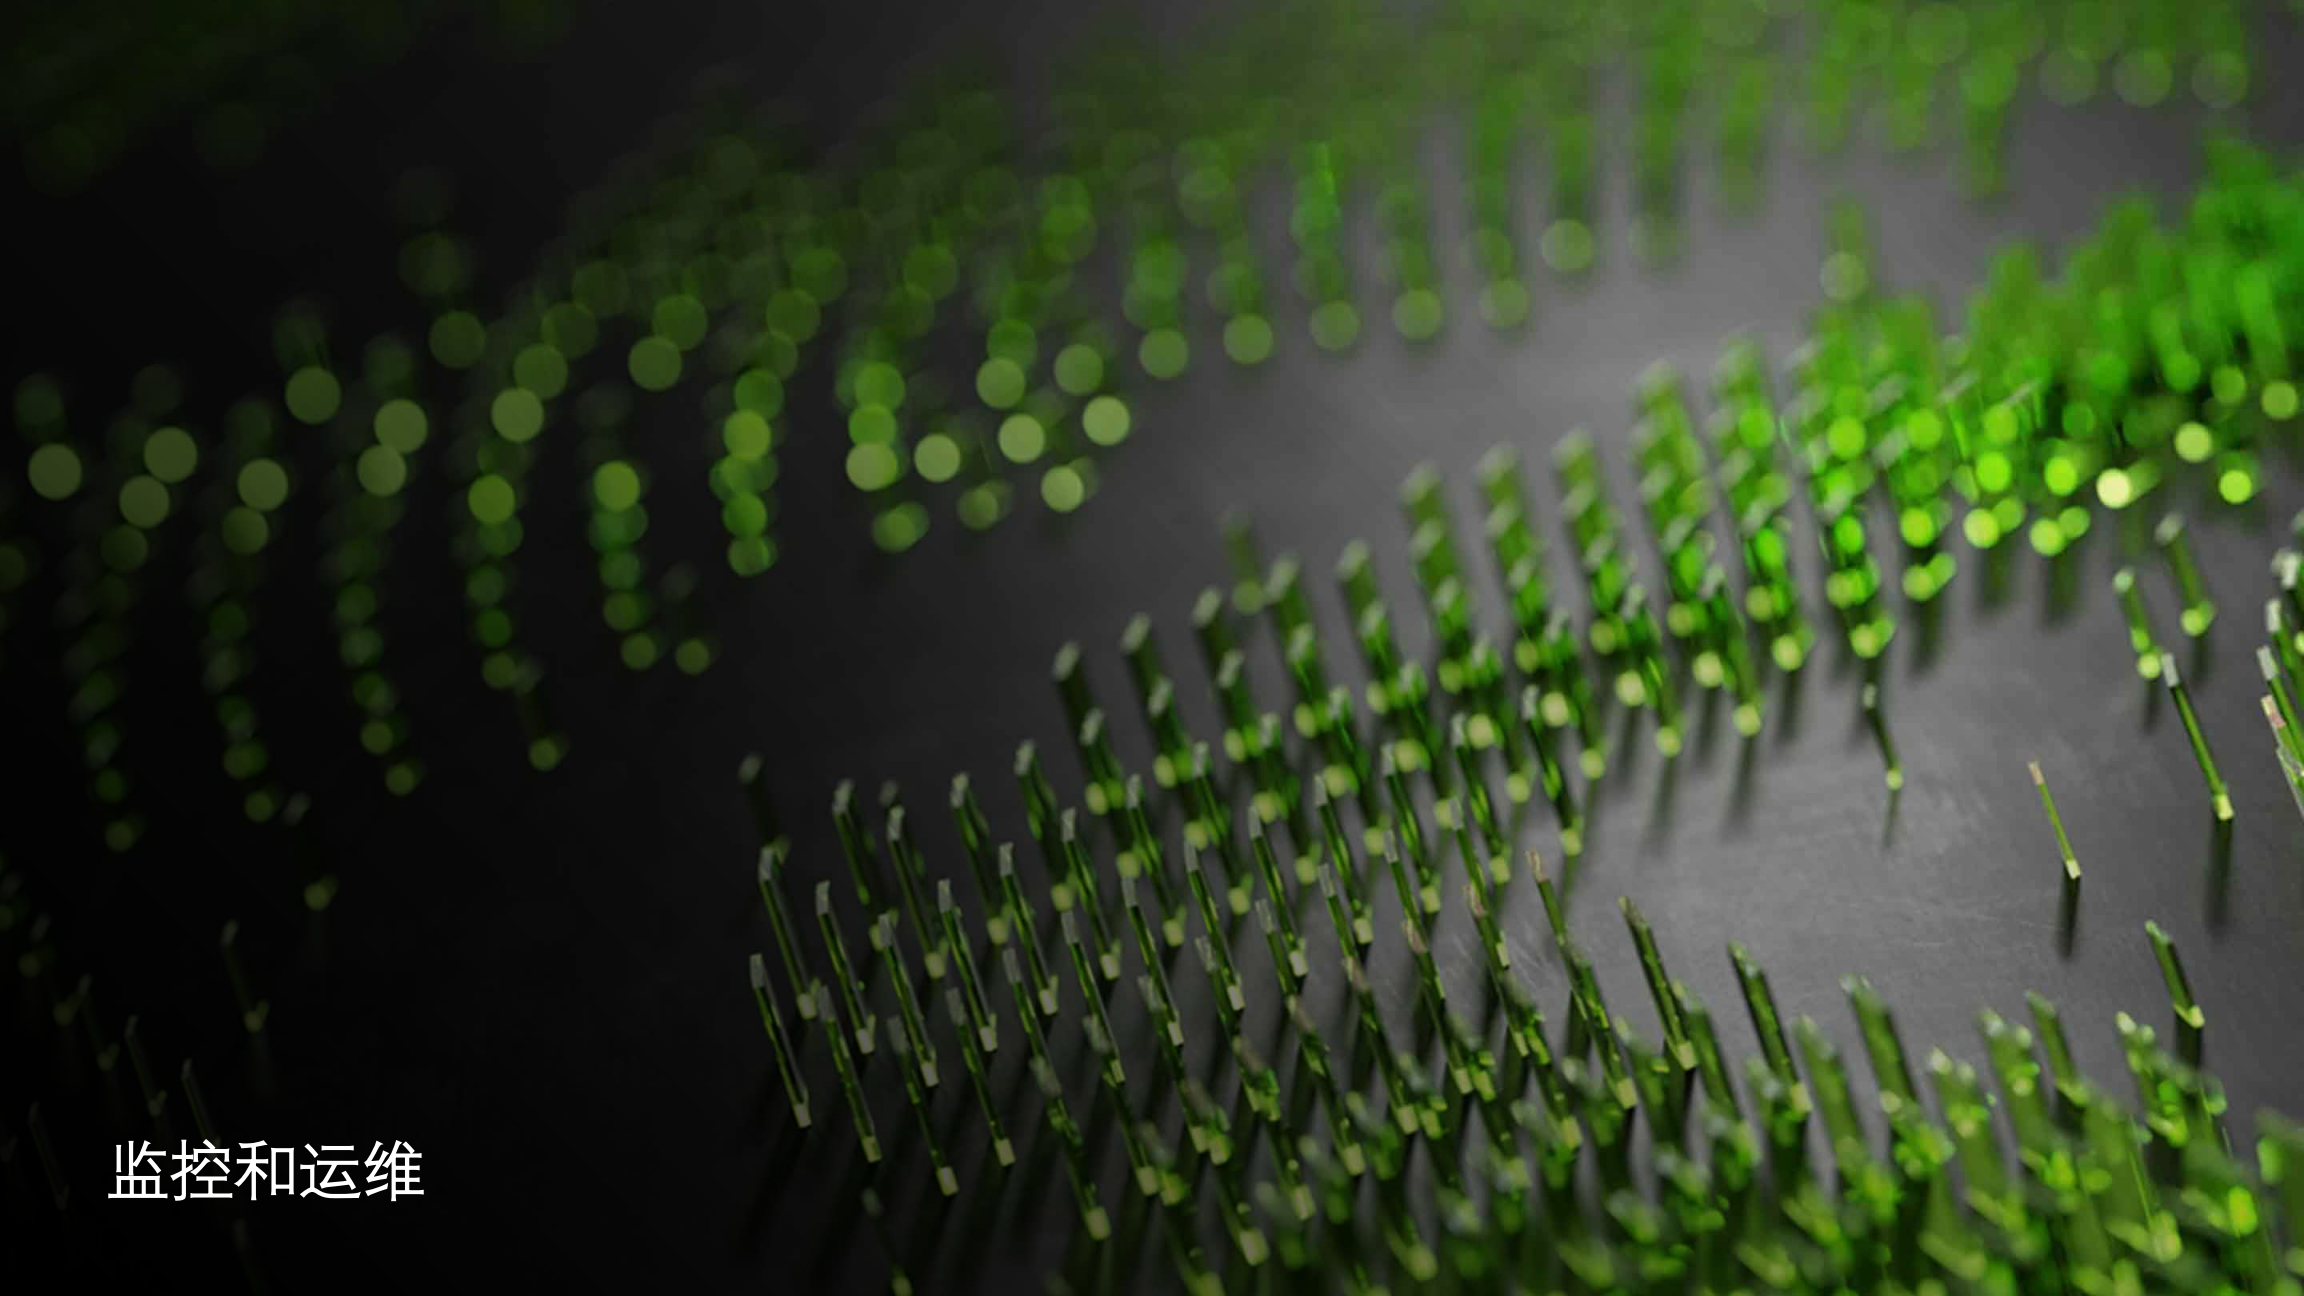
监控和运维

UFM PLATFORMS PORTFOLIO

UFM Telemetry
Real-Time Monitoring

UFM Enterprise
Management, Monitoring & Orchestration

(UFM Enterprise includes UFM Telemetry)

UFM Cyber-AI
Cyber Intelligence and Analytics

(UFM Cyber-AI includes UFM Enterprise)

UFM IN THE FABRIC

- Manages Subnet Manager and Sharp Services
- Software or appliance form factor
- High availability - 2 or more
- Switch and adapter management
- Full management or monitoring only
- Layer 2 level monitoring
- REST API for configuration/monitoring
- Single Interface for all network

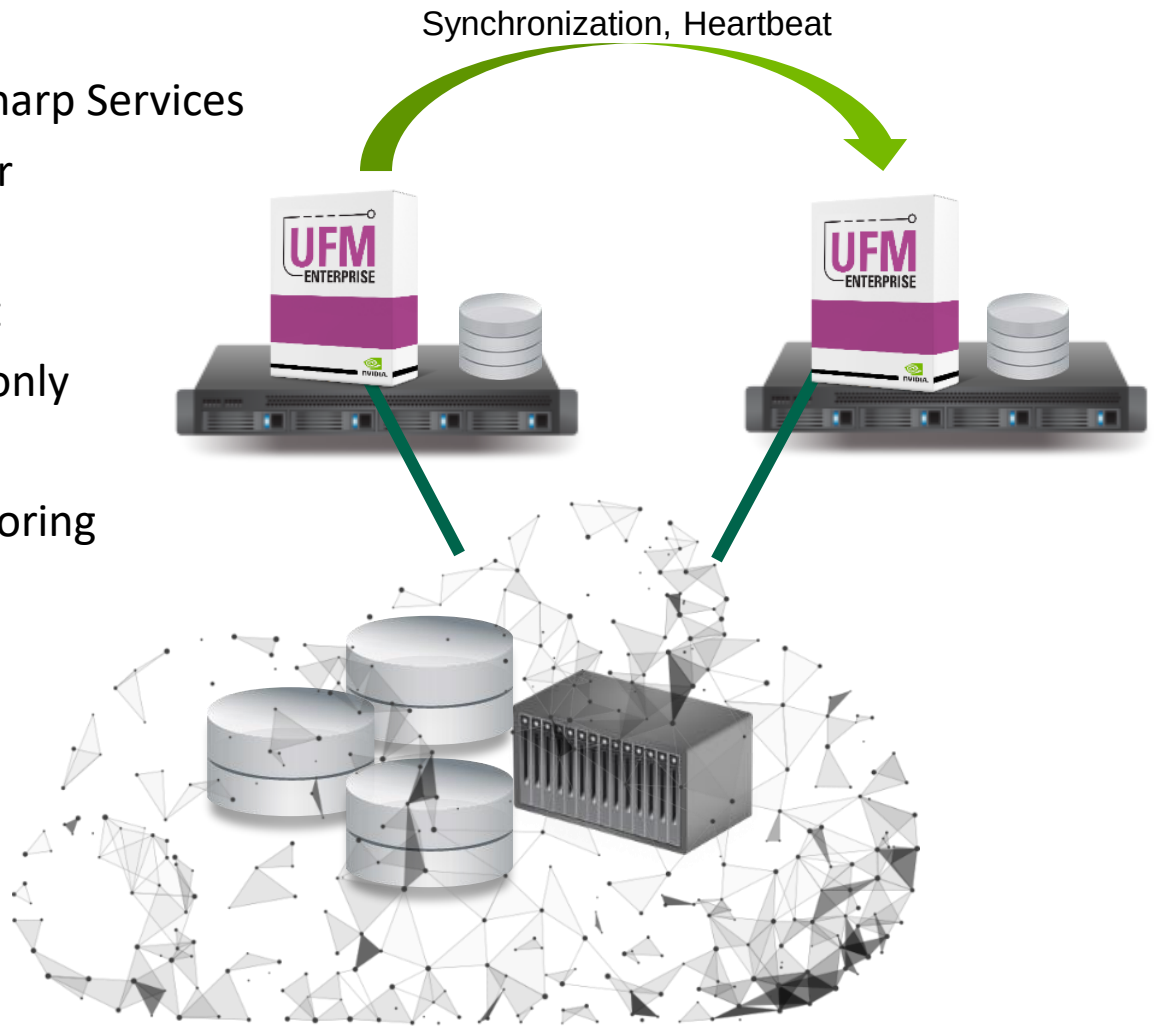

CENTRALIZE DEVICE MANAGEMENT

- Manage an inventory of assets, switches and nodes
- Centrally upgrade firmware and software across all managed and unmanaged systems

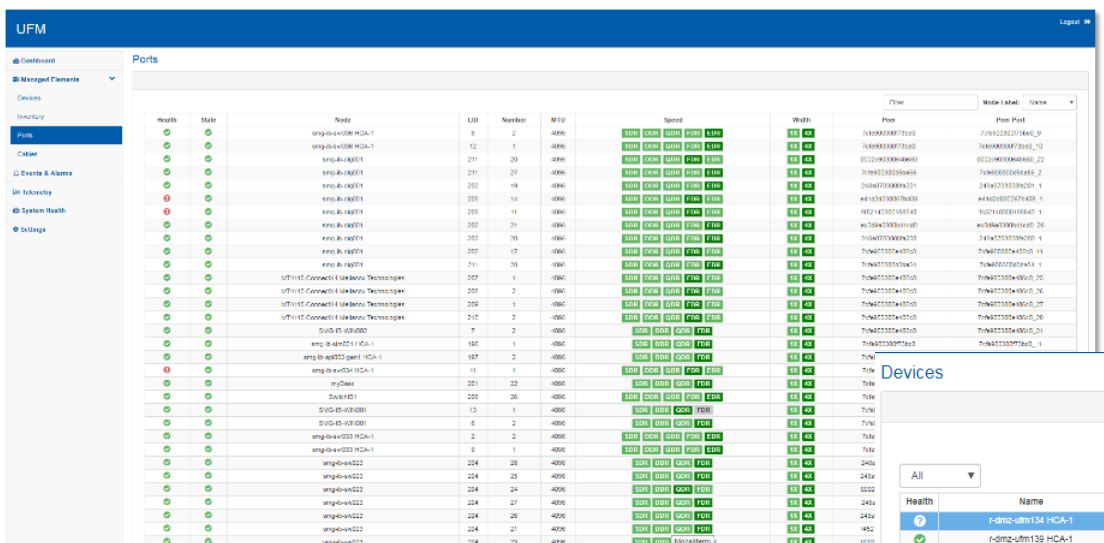

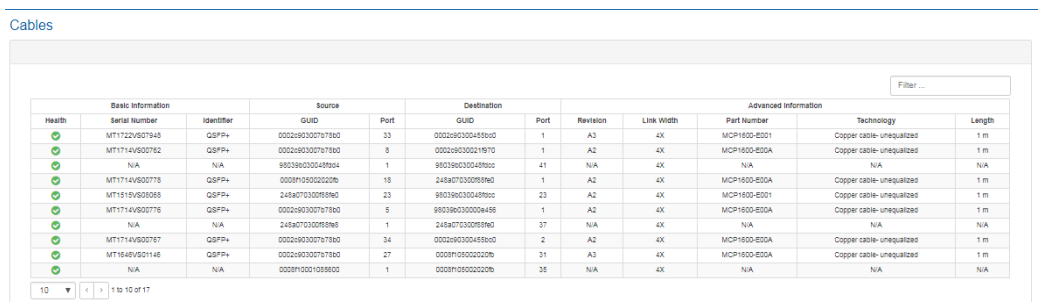

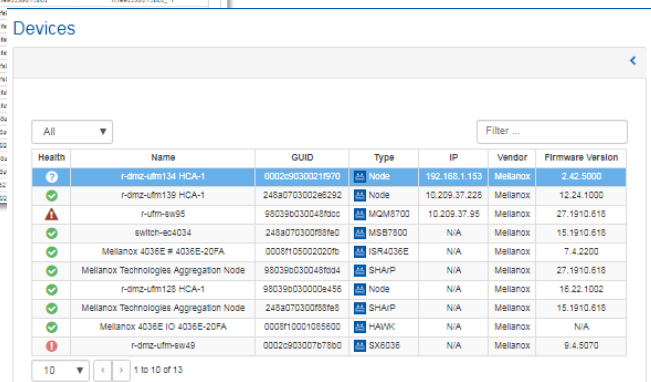

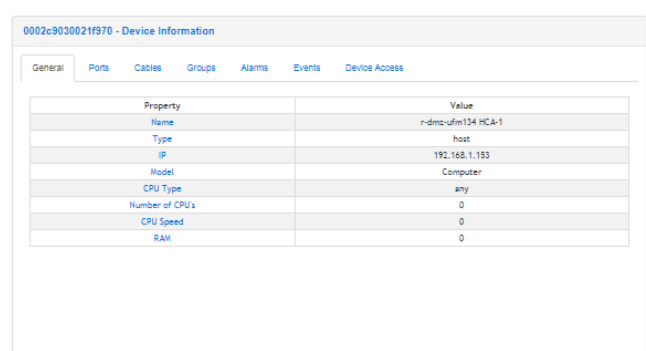

UFM DASHBOARD

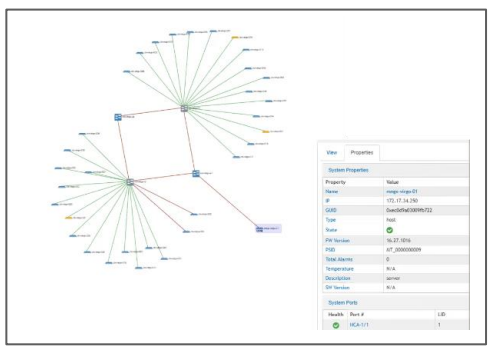

Network Validation

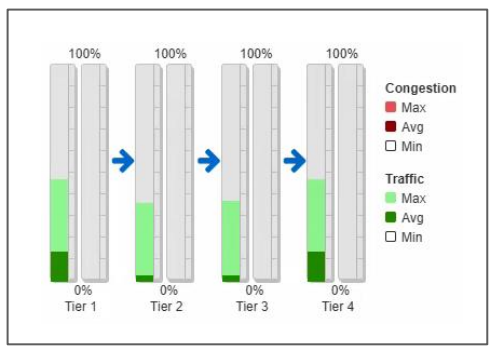

Congestion Mapping

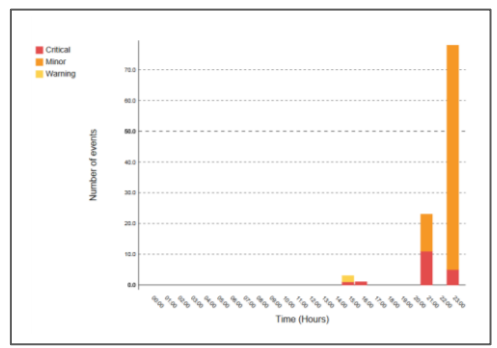

Health Reports

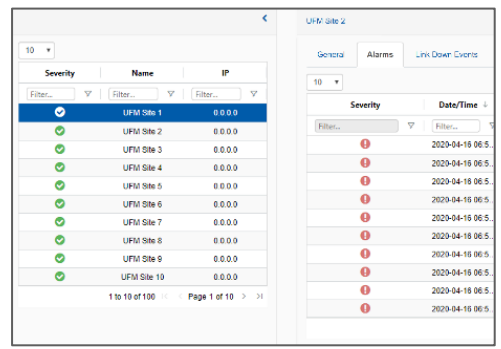

Inventory Mapping

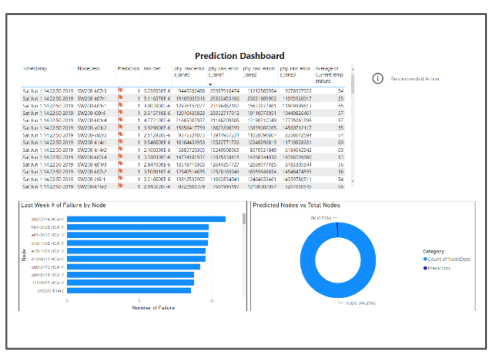

Prediction Dashboard

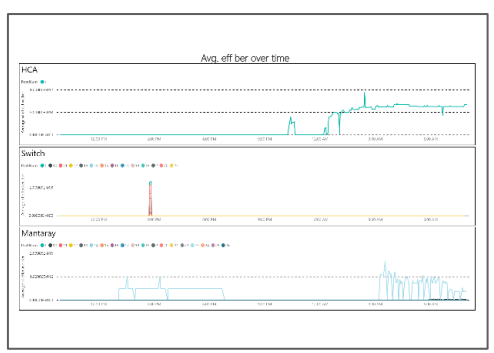

Real-Time Analysis

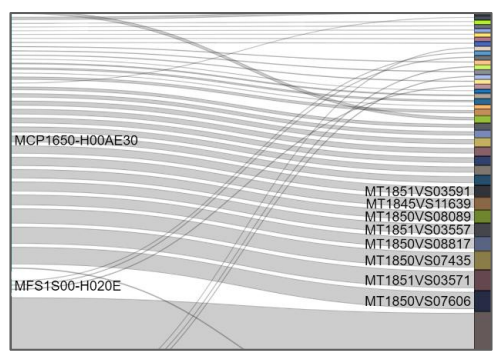

Performance Monitoring

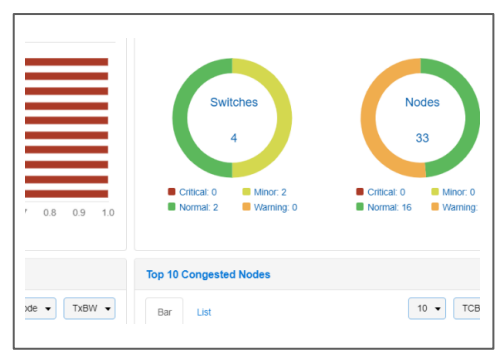

Secure Cable Management

CENTRALIZE DEVICE MANAGEMENT

- Manage an inventory of assets, switches and nodes
- Centrally upgrade firmware and software across all managed and unmanaged systems

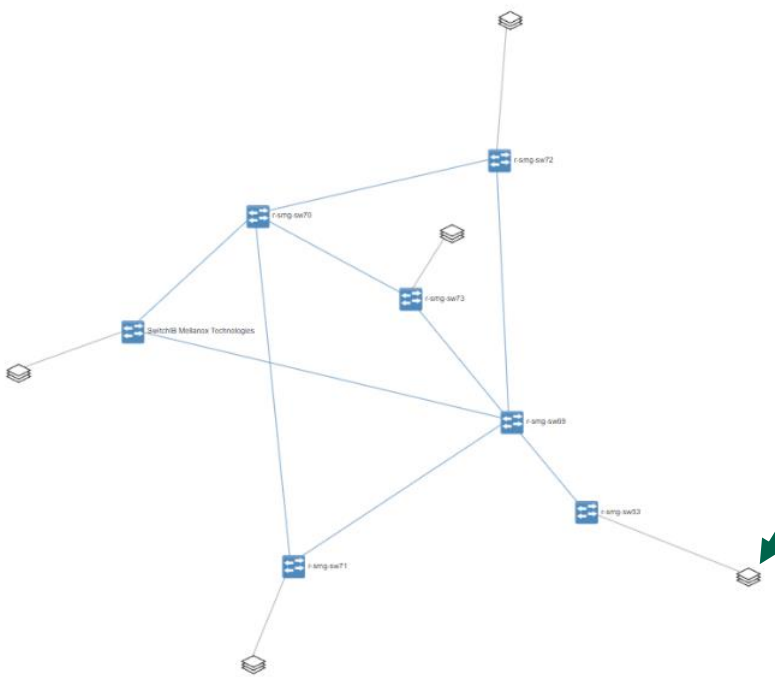

Groups

+ New

Groups

All | Filter ... | CSV

Severity	Name	Description	Type
✓	Suppressed_Devices	No event notifications issued	General
✓	Alarmed_Devices	Devices with alarms	General
✓	Rack_1	N/A	Rack

10 | 1 to 3 of 3

Rack_1 - Members

Filter ... | CSV

Name	GUID	IP
R-SMG161	e41d2d0300a20c0c	0.0.0.0
R-SMG165	e41d2d030067b0e8	0.0.0.0

10 | 1 to 2 of 2

MULTI CLOUD SOLUTION

- Correlation between Pkey creation/GUID assignment and traffic utilization/congestion
- Single main dashboard for all managed cloud/clusters
- Alerts, Traffic utilization, Congestion, Cable Info, Health

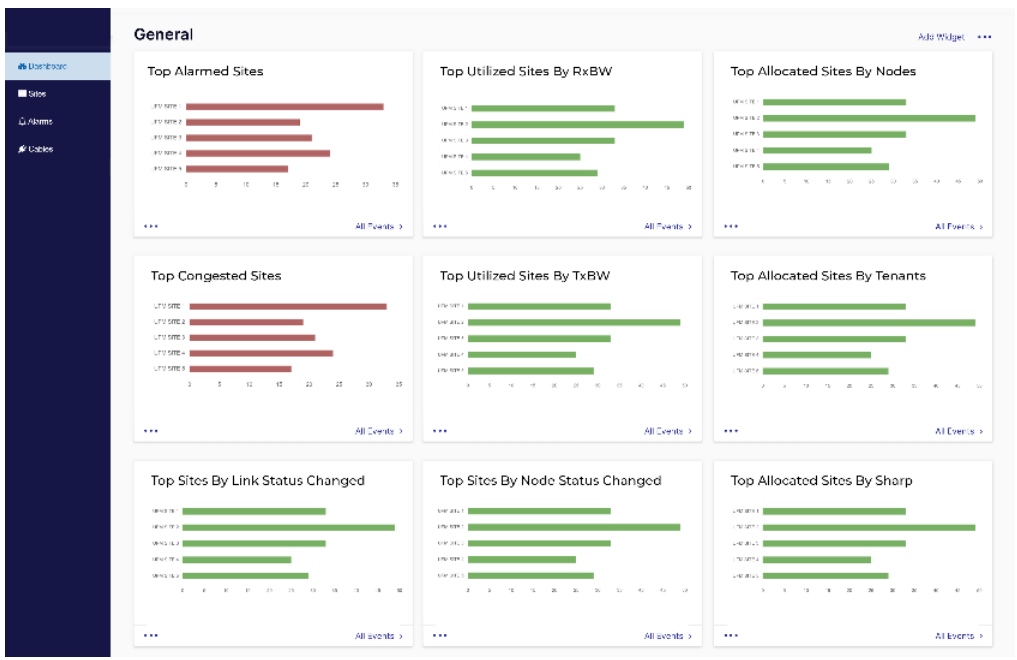

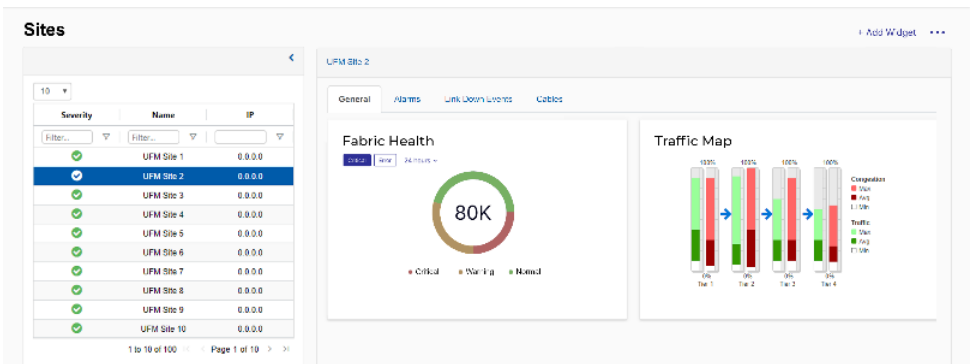

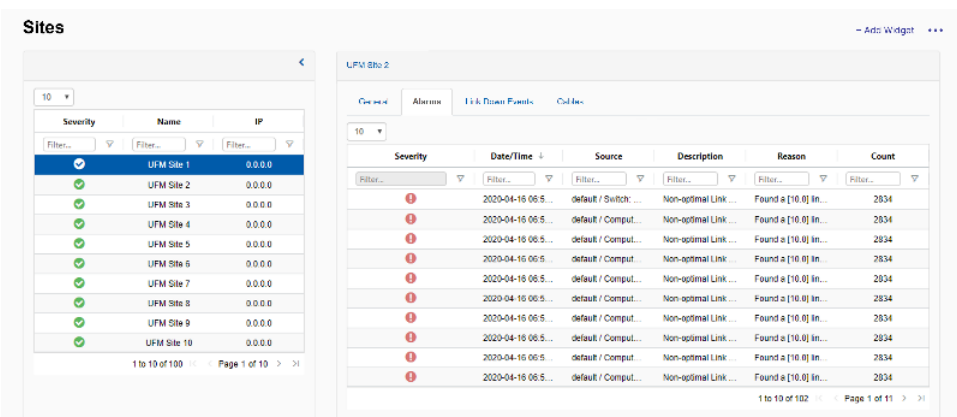

UFM IN CLOUD

Day 1 operations	Day 2 Operations
Fabric Bring up validation	Network Auto Provisioning
Cable check	Tenant Security and Isolation
Link check	Chassis Fault Detection
Connectivity Check	Network Congestion
BW Check	Network Issues
Latency Check	Network Analysis and Monitoring
Chassis Check	HA service for the network
Inventory Discovery + Health	Events and Alarms

UFM SECURITY FEATURES FOR CLOUD

Subject
Alert and action on SA_Key violation (detecting malicious queries and reporting)
Alert and action on SA DoSc (detecting and reporting attack)
Randomize SA_Key on SM start (in order to protect the SM from item 1)
Support for M_Key per port in SM and tools (protect fabric from malicious configuration)
Flows to isolate violator of security alerts (fabric operator action)
ConnectX-5 And 6 SLID anti-masquerading feature in steering logic
Prevent dDOS by malicious registration in SM
Secured Cable Management (Detecting cable changes)
Switch Port Bad Pkey Alert

UFM SECURITY EVENTS

ID	Subject	
256	Bad M_Key	Found bad Management key. Check your HCA driver or partition settings. Management Key: Enforces the control of a master subnet manager
257	Bad P_Key	Found a bad Partition key. Check your partitioning settings. Partition Key: Enforces membership. Administered through the subnet manager by the partition manager (PM).
258	Bad Q_Key	Found bad Queue key. Security error. Queue Key: Enforces access rights for reliable and unreliable datagram service (RAW datagram service type not included)
259	Bad P_Key Switch External Port	Found a bad Partition key. Check your partitioning settings. Partition Key: Enforces membership. Administered through the subnet manager by the partition manager (PM)
560	User Connected	User Connected
561	User Disconnected	User Disconnected
1300	SA Key violation	SA Key Volation Committed
1301	SGID Spoofed	SGID spoofed by VPort/port
1302	SA High Rate detected	Rate Limit Exceeded

UFM SOFTWARE ARCHITECTURE

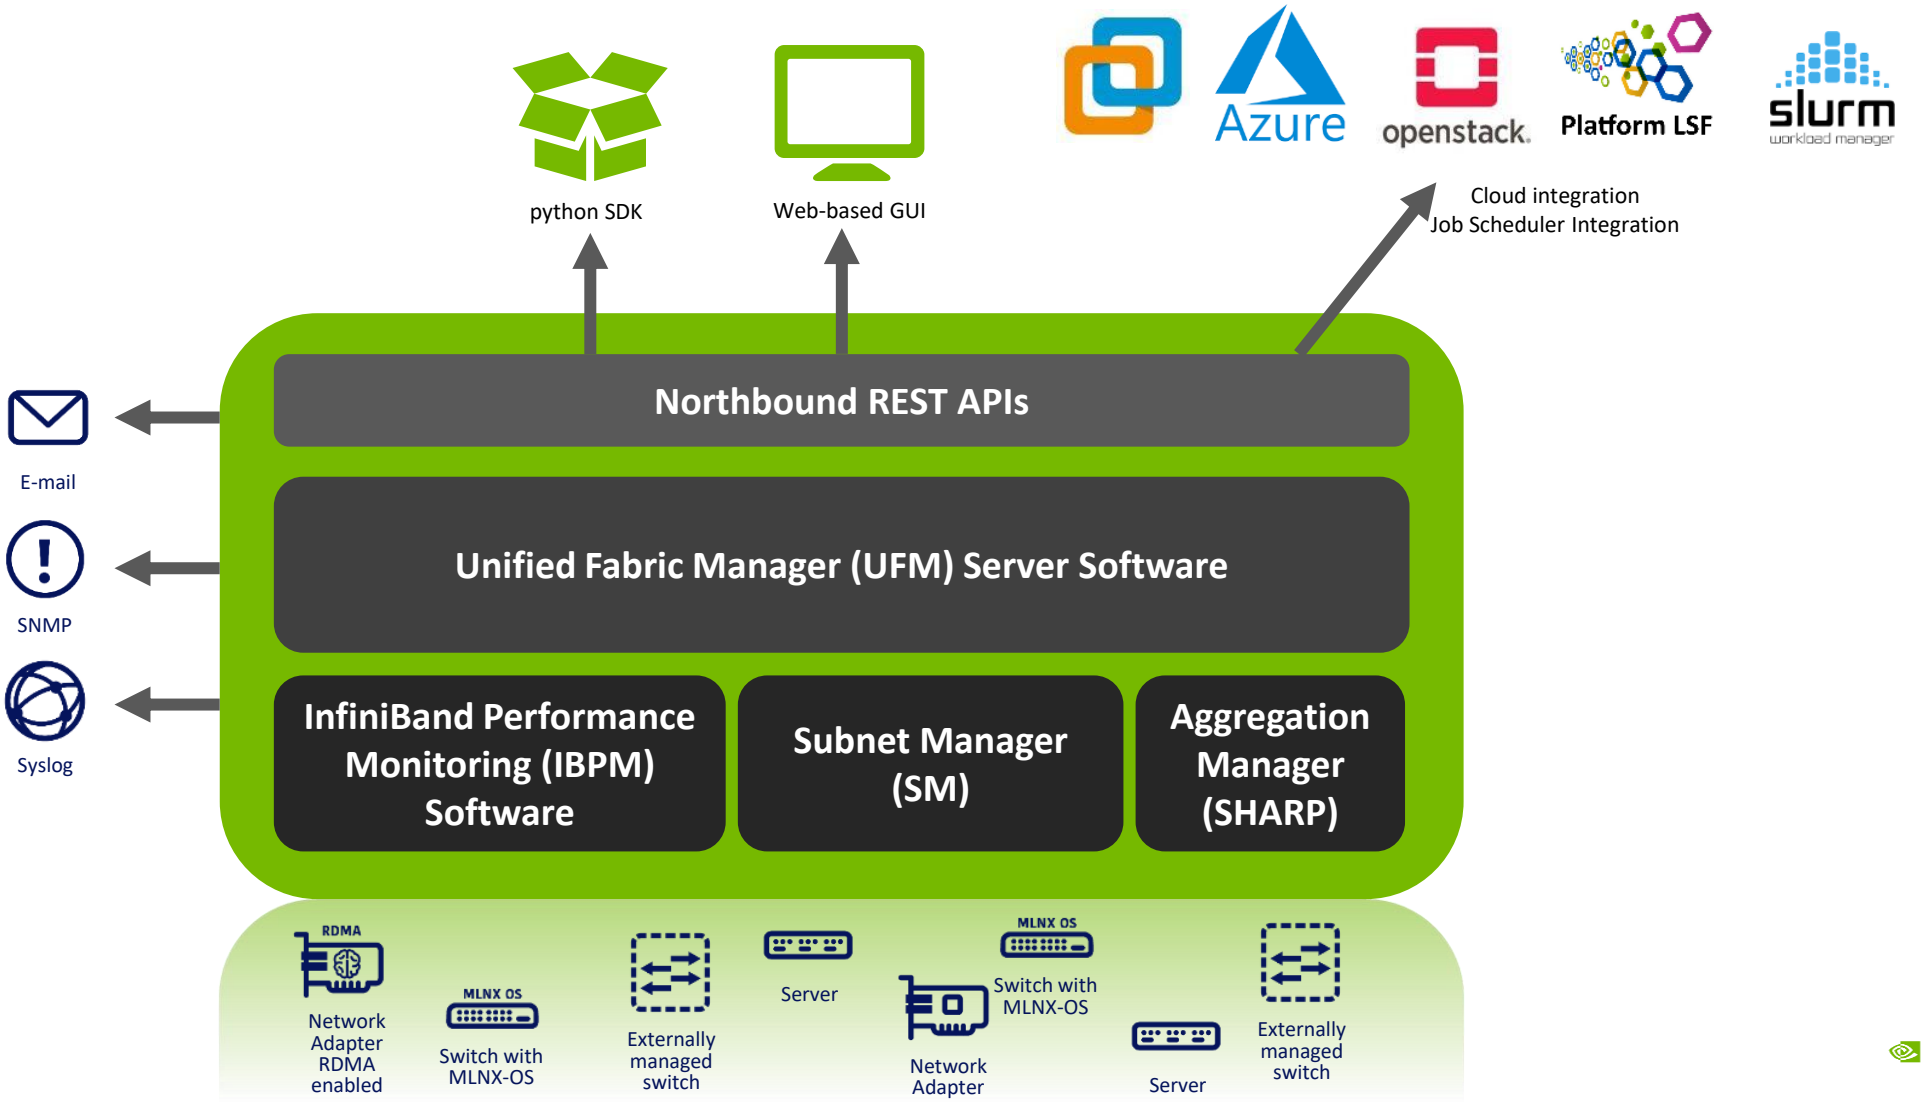

NVIDIA

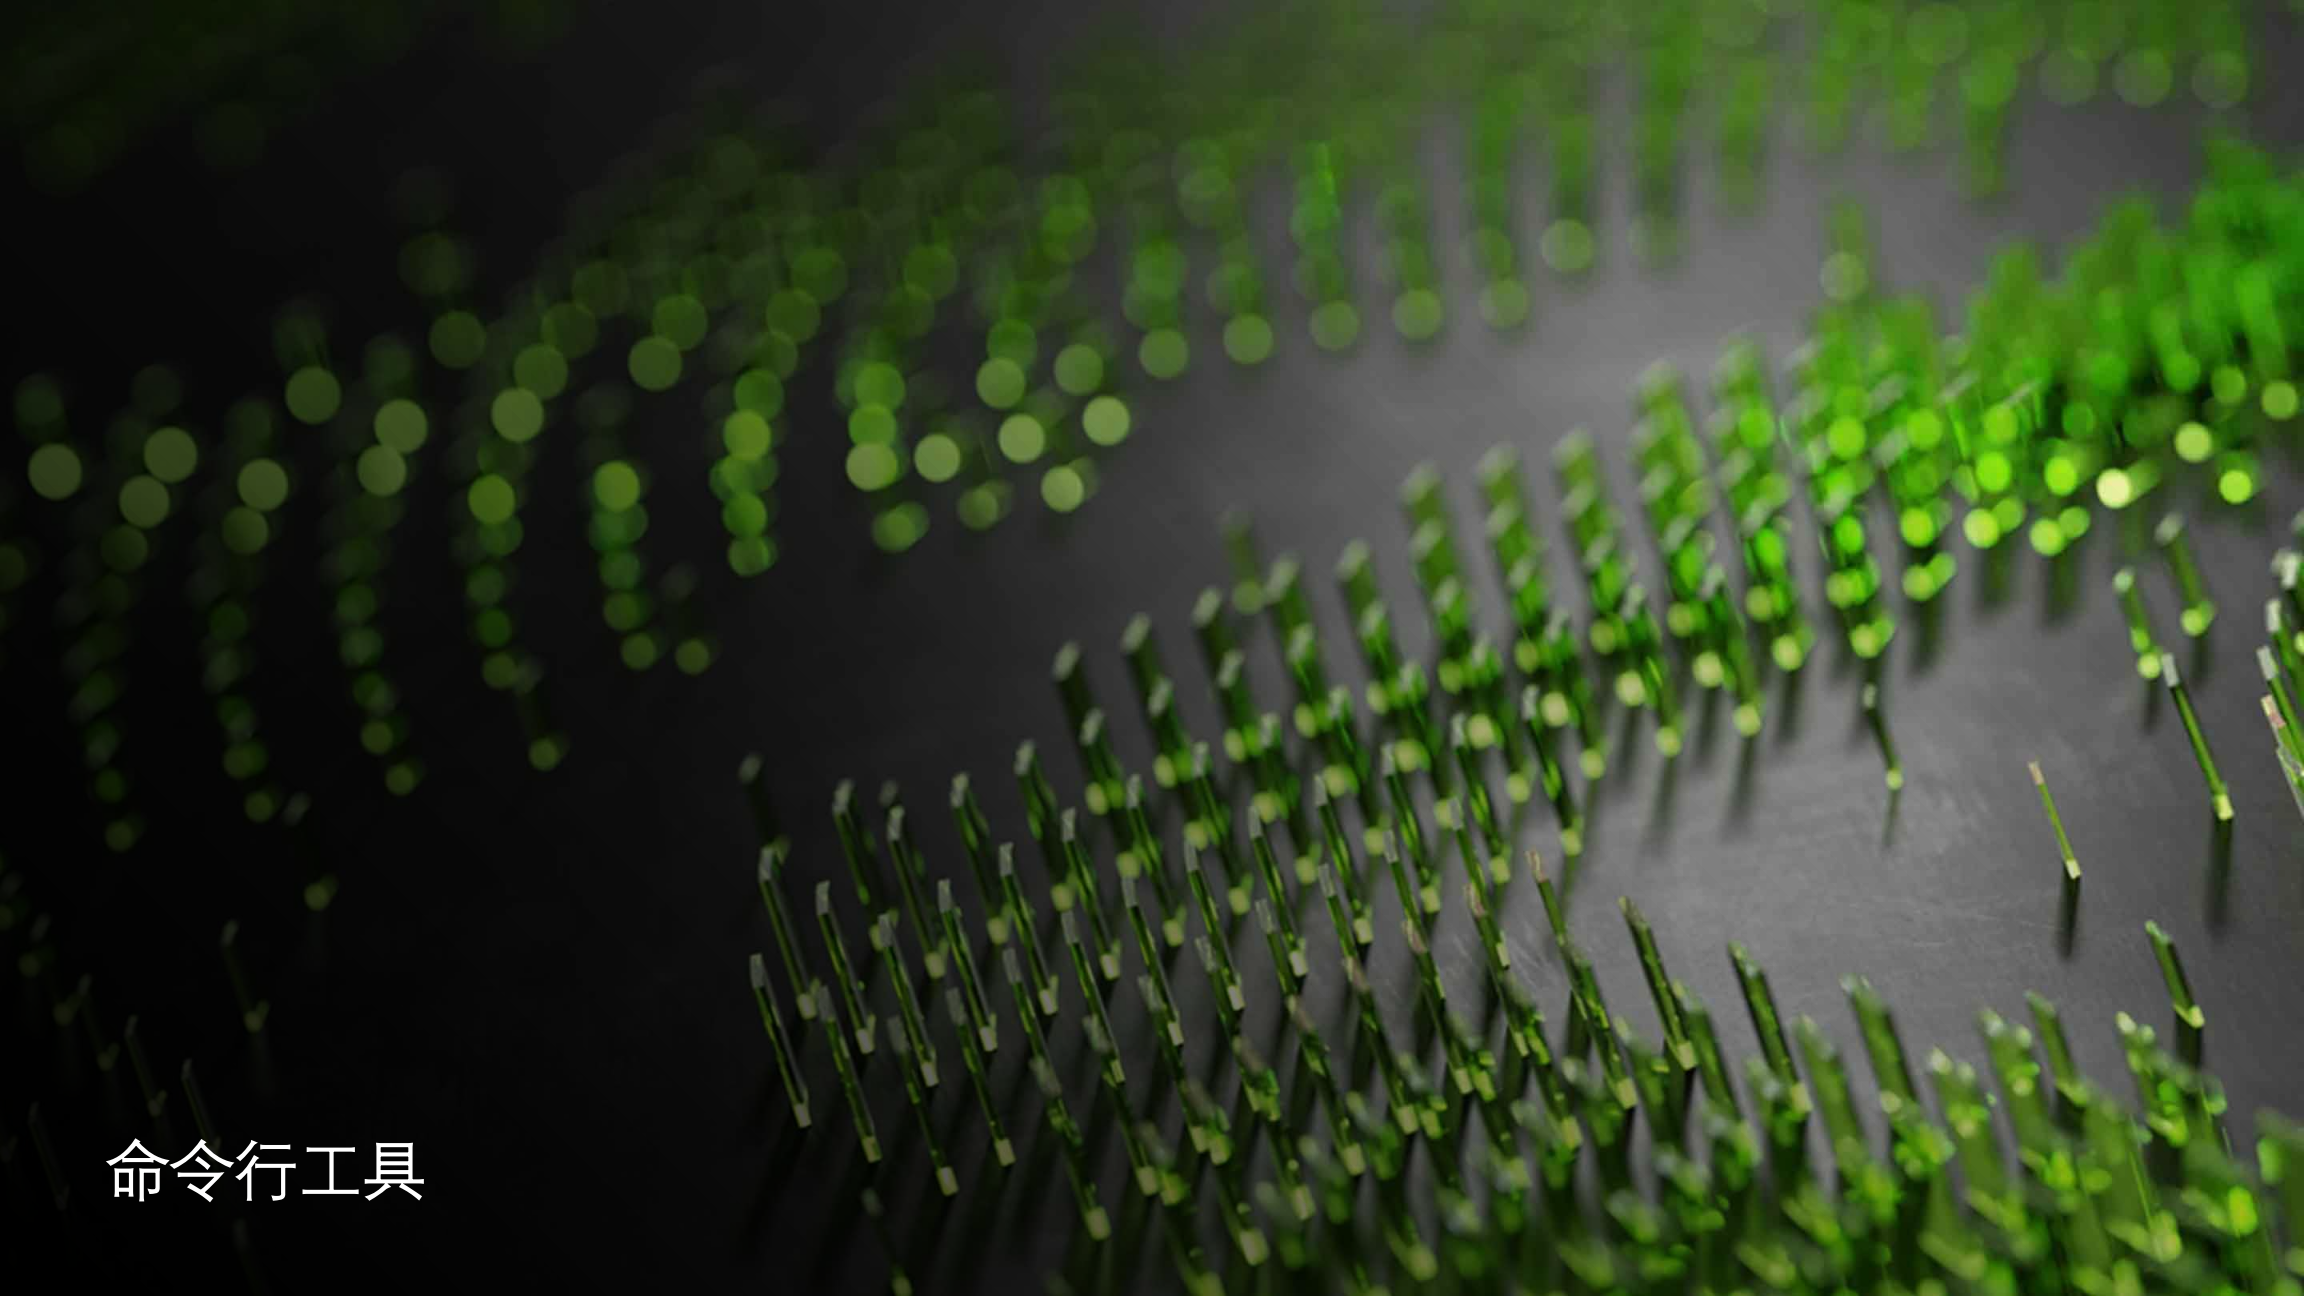
命令行工具

IBDIAGNET

Ibutils2

Scans the fabric using directed route packets and extracts all the available information regarding its connectivity and devices. An ibdiagnet run performs the following stages:

- Fabric discovery
- Duplicated GUIDs detection
- Links in INIT state and unresponsive links detection
- Counters fetch
- Error counters check
- Routing checks
- Link width and speed checks
- Alias GUIDs check
- Subnet Manager check
- Partition keys check

IBDIAGNET COMMAND

INSTALLATION

---

mkdir /tmp/mlnx

cd /tmp/mlnx

cp <path>/ibdiagnet_monitor_4.5.tgz .

tar zxf ibdiagnet_monitor_4.5.tgz

ibdiagnet_monitor_4.5.tgz

export IBDIAGNET_PLUGINS_PATH=/tmp/mlnx/usr/share/ibdiagnet2.1.1/plugins

export LD_LIBRARY_PATH=/tmp/mlnx/usr/lib

/tmp/mlnx/usr/bin/ibdiagnet <relevent flags>

/tmp/mlnx/usr/bin/ibdiagnet -pc --pm_pause_time 300 -P all=1 --get_cable_info --get_phy_info

Logs will be at the same place in /var/tmp/ibdiagnet2/*

#iblinkinfo

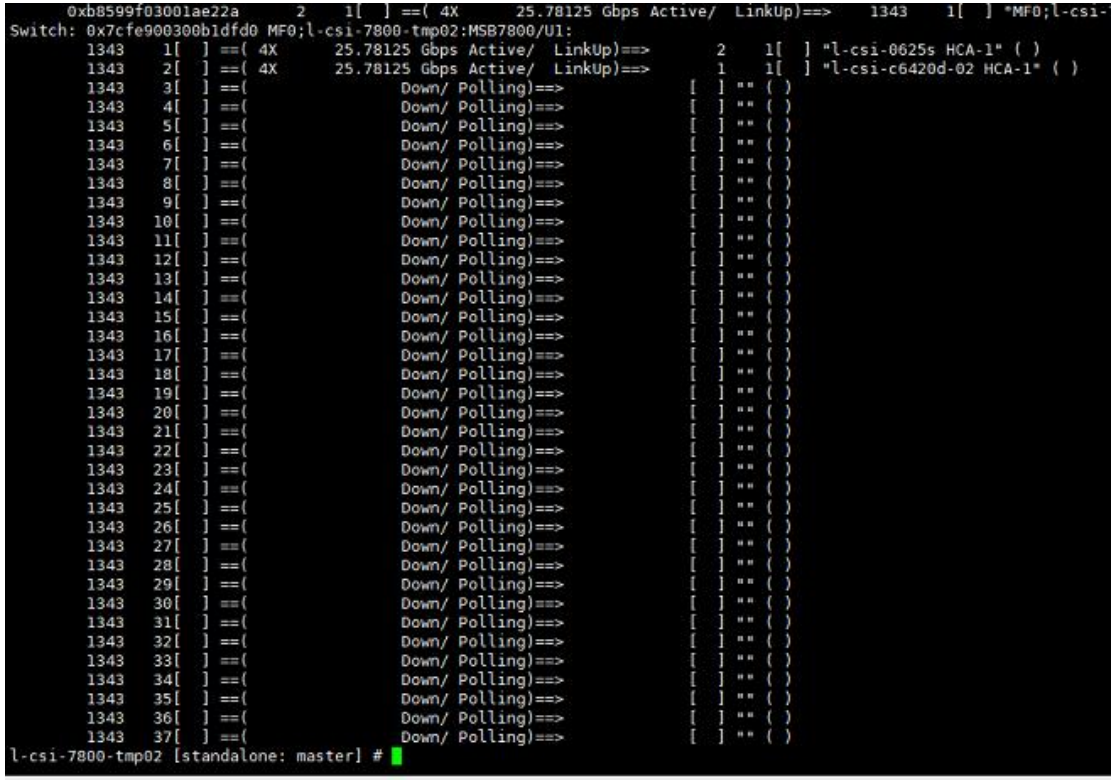
0xb8599f03001ae22a 2 1[ ] ==(4X 25.78125 Gbps Active/ LinkUp)==> 1343 1[ ] "MF0;l-csi-
Switch: 0x7cfe900300b1dfd0 MF0;l-csi-7800-tmp02:MSB7800/U1:
1343 1[ ] ==(4X 25.78125 Gbps Active/ LinkUp)==> 2 1[ ] "l-csi-0625s HCA-1" ( )
1343 2[ ] ==(4X 25.78125 Gbps Active/ LinkUp)==> 1 1[ ] "l-csi-c6420d-02 HCA-1" ( )
1343 3[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 4[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 5[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 6[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 7[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 8[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 9[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 10[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 11[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 12[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 13[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 14[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 15[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 16[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 17[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 18[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 19[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 20[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 21[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 22[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 23[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 24[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 25[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 26[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 27[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 28[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 29[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 30[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 31[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 32[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 33[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 34[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 35[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 36[ ] ==(Down/ Polling)==> [ ] "" ( )
1343 37[ ] ==(Down/ Polling)==> [ ] "" ( )
l-csi-7800-tmp02 [standalone: master] #

#ibnetdiscover

```
l-csi-7800-tmp02 [standalone: master] # fae ibnetdiscover
#
# Topology file: generated on Sun Mar 15 10:09:53 2020
#
# Initiated from node 7cfe900300b1dfd0 port 7cfe900300b1dfd0

vendid=0x2c9
devid=0xcf08
sysimgguid=0x7cfe900300b1dfd0
switchguid=0x7cfe900300b1dfd0(7cfe900300b1dfd0)
Switch  37 "S-7cfe900300b1dfd0"         # "MF0;l-csi-7800-tmp02:MSB7800/U1" enhanced port 0 lid 1343 lmc 0
[1]     "H-b8599f03001ae22a"[1](b8599f03001ae22a)              # "l-csi-0625s HCA-1" lid 2 4xEDR
[2]     "H-b8599f0300fe4d50"[1](b8599f0300fe4d50)              # "l-csi-c6420d-02 HCA-1" lid 1 4xEDR

vendid=0x2c9
devid=0x1019
sysimgguid=0xb8599f0300fe4d50
caguid=0xb8599f0300fe4d50
Ca      1 "H-b8599f0300fe4d50"          # "l-csi-c6420d-02 HCA-1"
[1](b8599f0300fe4d50)   "S-7cfe900300b1dfd0"[2]         # lid 1 lmc 0 "MF0;l-csi-7800-tmp02:MSB7800/U1" lid 1343 4xEDR

vendid=0x2c9
devid=0x1019
sysimgguid=0xb8599f03001ae22a
caguid=0xb8599f03001ae22a
Ca      1 "H-b8599f03001ae22a"          # "l-csi-0625s HCA-1"
[1](b8599f03001ae22a)   "S-7cfe900300b1dfd0"[1]         # lid 2 lmc 0 "MF0;l-csi-7800-tmp02:MSB7800/U1" lid 1343 4xEDR
l-csi-7800-tmp02 [standalone: master] #
```

#ibstat

```
[root@l-csi-0636s ~]# ibstat
CA 'mlx5_0'
        CA type: MT4119
        Number of ports: 1
        Firmware version: 16.26.1040
        Hardware version: 0
        Node GUID: 0xec0d9a0300ced24a
        System image GUID: 0xec0d9a0300ced24a
        Port 1:
                State: Down
                Physical state: Polling
                Rate: 10
                Base lid: 65535
                LMC: 0
                SM lid: 0
                Capability mask: 0x2651e848
                Port GUID: 0xec0d9a0300ced24a
                Link layer: InfiniBand
CA 'mlx5_1'
        CA type: MT4119
        Number of ports: 1
        Firmware version: 16.26.1040
        Hardware version: 0
        Node GUID: 0xec0d9a0300ced24b
        System image GUID: 0xec0d9a0300ced24a
        Port 1:
                State: Down
                Physical state: Disabled
                Rate: 10
                Base lid: 65535
                LMC: 0
                SM lid: 0
                Capability mask: 0x2651e848
                Port GUID: 0xec0d9a0300ced24b
                Link layer: InfiniBand
[root@l-csi-0636s ~]#
```

#ibv_devinfo

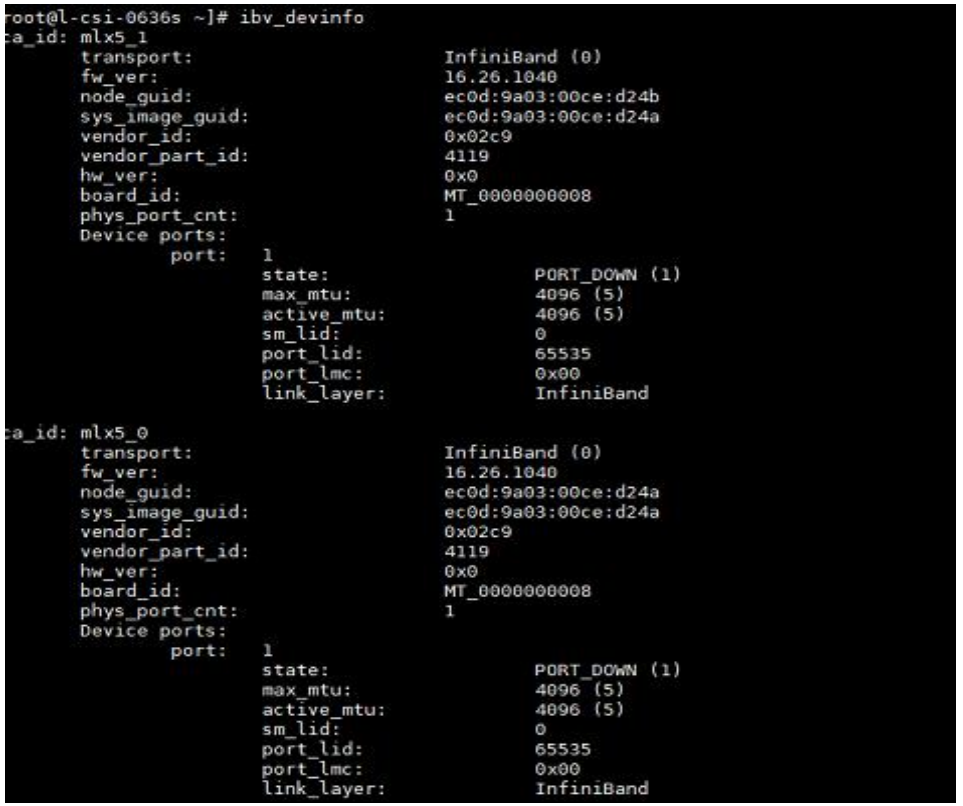
root@l-csi-0636s ~]# ibv_devinfo
ca_id: mlx5_1
transport: InfiniBand (0)
fw_ver: 16.26.1040
node_guid: ec0d:9a03:00ce:d24b
sys_image_guid: ec0d:9a03:00ce:d24a
vendor_id: 0x02c9
vendor_part_id: 4119
hw_ver: 0x0
board_id: MT_0000000008
phys_port_cnt: 1
Device ports:
port: 1
state: PORT_DOWN (1)
max_mtu: 4096 (5)
active_mtu: 4096 (5)
sm_lid: 0
port_lid: 65535
port_lmc: 0x00
link_layer: InfiniBand
ca_id: mlx5_0
transport: InfiniBand (0)
fw_ver: 16.26.1040
node_guid: ec0d:9a03:00ce:d24a
sys_image_guid: ec0d:9a03:00ce:d24a
vendor_id: 0x02c9
vendor_part_id: 4119
hw_ver: 0x0
board_id: MT_0000000008
phys_port_cnt: 1
Device ports:
port: 1
state: PORT_DOWN (1)
max_mtu: 4096 (5)
active_mtu: 4096 (5)
sm_lid: 0
port_lid: 65535
port_lmc: 0x00
link_layer: InfiniBand

#smpquery

```
[root@l-csi-0636s ~]# smpquery -h

Usage: smpquery [options] <op> <dest dr_path|lid|guid> [op params]

Supported ops (and aliases, case insensitive):
  NodeInfo (NI) <addr>
  NodeDesc (ND) <addr>
  PortInfo (PI) <addr> [<portnum>]
  PortInfoExtended (PIE) <addr> [<portnum>]
  SwitchInfo (SI) <addr>
  PKeyTable (PKeys) <addr> [<portnum>]
  SL2VLTable (SL2VL) <addr> [<portnum>]
  VLArbitration (VLArb) <addr> [<portnum>]
  GUIDInfo (GI) <addr>
  MlnxExtPortInfo (MEPI) <addr> [<portnum>]


Options:
  --combined, -c          use Combined route address argument
  --node-name-map <file>  node name map file
  --extended, -x          use extended speeds
  --config, -z <config>   use config file, default: /etc/infiniband-diags/ibdiag.conf
  --Ca, -C <ca>           Ca name to use
  --Port, -P <port>       Ca port number to use
  --Direct, -D            use Direct address argument
  --Lid, -L               use LID address argument
  --Guid, -G              use GUID address argument
  --timeout, -t <ms>      timeout in ms
  --sm_port, -s <lid>     SM port lid
  --show_keys, -K         display security keys in output
  --m_key, -y <key>       M_Key to use in request
  --errors, -e            show send and receive errors
  --verbose, -v           increase verbosity level
  --debug, -d             raise debug level
  --help, -h              help message
  --version, -V           show version
  --m_key_files, -w <dir_path>  Path to direcory that include m_key files

Examples:
  smpquery portinfo 3 1                       # portinfo by lid, with port modifier
  smpquery -G switchinfo 0x2C9000100D051 1    # switchinfo by guid
  smpquery -D nodeinfo 0                             # nodeinfo by direct route
  smpquery -c nodeinfo 6 0,12                 # nodeinfo by combined route

[root@l-csi-0636s ~]#
```

OTHER COMMAND

- #ibswitches
- #ibnodes
- #ibhosts
- #ofed_info –S
- #mlxlink –d lid-<lid#> -p <port#>
- #perfquery
- #ibportstate
- #flint
- #mlxvpd
- #mlxup
- #mlxconfig
- #mlxfwmanage
- #ib_write_bw
- #ib_write_lat
- #ibping
- #ib_read_bw
- #ib_read_lat

HOST TOOLS

Sysinfo-snapshot

Windows

<installation_directory>\ManagementTools\MLNX_System_Snapshot.exe

Linux

#./sysinfo-snapshot.py

ESXi

esxi-sysinfo-snapshot.py


NVIDIA