

# RoCE & QoS

Jan 05, 2015

V-1.0.0

Changqing Li  
licq@mellanox.com



- Background
- What is RoCE
- RoCE Packet Format
- RoCE packet encapsulation
- RoCE QoS with Verbs API
- RoCE QoS with RDMA\_CM API
- Differences between RDMA programming over IB and RoCE
- Performance
- RoCEv2
- Reference
- Q&A

- Advantage of RDMA(Zero-copy, full protocol offload, high performance, good latency, low CPU utilization).
- Lots of Ethernet device and equipment deployed in data center
- Legacy TCP/IP stack becoming a problem with the growth of BW
- IEEE 802.1 DCB ensures lossless Ethernet network.

## ■ RoCE - RDMA over Converged Ethernet

- 802.1Qbb - PFC
- 802.1az - ETS
- 802.1AB - LLDP, DCBx
- 802.1Qau - CN


## ■ Defined by IBTA

- Annex A16: RDMA over Converged Ethernet (RoCE) of “InfiniBand Architecture Specification Volume 1 Release 1.2.1”

## Legacy RoCE Packet



## RoCE Packet with Vlan Tag



- ⊕ Frame 10100: 94 bytes on wire (752 bits), 94 bytes captured (752 bits)
- ⊖ Ethernet II, Src: Mellanox\_47:66:80 (f4:52:14:47:66:80), Dst: Mellanox\_2e:71:50 (f4:52:14:2e:71:50)
  - ⊕ Destination: Mellanox\_2e:71:50 (f4:52:14:2e:71:50)
  - ⊕ Source: Mellanox\_47:66:80 (f4:52:14:47:66:80)  
Type: 802.1Q Virtual LAN (0x8100)
- ⊖ 802.1Q Virtual LAN, PRI: 0, CFI: 0, ID: 100
  - 000. .... = Priority: Best Effort (default) (0)
  - ...0 .... = CFI: Canonical (0)
  - .... 0000 0110 0100 = ID: 100
  - Type: RDMA over Converged Ethernet (0x8915)
- ⊖ InfiniBand
  - ⊕ Global Route Header
  - ⊕ Base Transport Header
  - ⊖ RETH - RDMA Extended Transport Header
    - Virtual Address: 140341841166784
    - Remote Key: 3355510020
    - DMA Length: 3

## Source MAC Address

22:16	Eth/IB	mlid / smac_index	0	For Ethernet QP - index to MAC table this QP is connected to. For IB - LMC bits of LID
-------	--------	-------------------	---	--

## Destination MAC Address

024h	31:16		reserved		
	15:0	Eth	dmac[47:32]	1	Upper 16 bits of Remote MAC address. Valid for a QP connected to an Ethernet port
028h	31:0	Eth	dmac[31:0]	1	Upper 32 bits of Remote MAC address. Valid for a QP connected to an Ethernet port

## Vlan ID

22:16	Eth	vlan_index	0	Index to the VLAN-ID Table. This VLAN-ID will be inserted for outgoing packets, to be checked on received packets. Priority bits are according to the SL in the Schedule Queue number.
-------	-----	------------	---	--

## UP (Vlan Priority)

**Table 59 - sched\_queue Field Parameter Description**

Bits	Name	Description
6	port	Port index: 0x0 - port 1 0x1 - port 2
5:2	SL	IB Service Level
5:3	priority	Ethernet user Priority



## User space RoCE application needs to do:

```
int flags      = IBV_QP_STATE;
```

```
attr->qp_state = IBV_QPS_RTR;
```

```
attr->ah_attr.is_global      = 1;
```

```
attr->ah_attr.grh.dgid       = <Destination GID>;
```

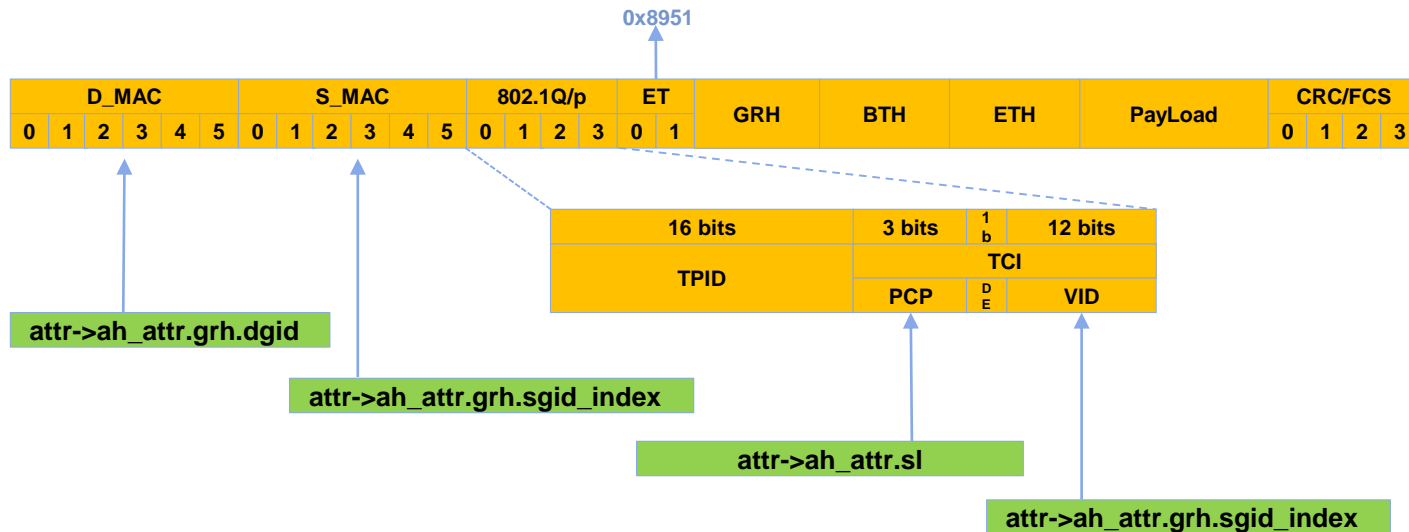
```
attr->ah_attr.grh.sgid_index = <Local GID index>;
```

```
attr->ah_attr.grh.hop_limit  = 1;
```

```
attr->ah_attr.sl             = <SL>;
```

```
ibv_modify_qp(qp, attr, flags);
```

# RoCE Packet Encapsulation – Hardware Behavior



# RoCE Packet Encapsulation – Driver Behavior



```
-> __uverbs_modify_qp
    -> ib_query_gid(sgid_index, &sgid)
    -> rdma_addr_find_dmac_by_grh(sgid, dgid, (char*)dmac, &vlan_id)
        -> sip = rdma_gid2ip(sgid)
        -> dip = rdma_gid2ip(dgid)
        -> rdma_resolve_ip(sip, dip, &dev_addr)
        -> memcpy(dmac, dev_addr.dst_dev_addr, ETH_ALEN);    // Got dmac
        -> dev = dev_get_by_index(bound_dev_if)
        -> vlan_id = rdma_vlan_dev_vlan_id(dev)                // Got Vlan ID

    -> rdma_addr_find_smac_by_sgid(sgid, (char*)smac)
        -> sip = rdma_gid2ip(sgid)
        -> rdma_translate_ip(sip, &dev_addr)
        -> memcpy(smac, dev_addr.src_dev_addr, ETH_ALEN);    // Got dmac

    -> ib_modify_qp
        -> mlx4_ib_modify_qp
            -> __mlx4_ib_modify_qp
                -> mlx4_set_path
                    -> path->sched_queue = (ah->sl & 7) << 3;    // Set Prio
```

## ■ What will happen without PFC?

```
[sincereli@test04 ~]$ /opt/mvapich2/gdr/2.0/gnu/bin/mpirun_rsh -np 2 test04
test05 MV2_USE_RoCE=1 MV2_USE_CUDA=1 MV2_USE_GPUDIRECT=1 MV2_USE_RDMA_CM=1
/opt/mvapich2/gdr/2.0/gnu/libexec/mvapich2/osu_bw -d cuda D D
# OSU MPI-CUDA Bandwidth Test
# Send Buffer on DEVICE (D) and Receive Buffer on DEVICE (D)
# Size          Bandwidth (MB/s)
1                0.68
2                1.37
4                2.75
8                5.49
16               11.00
32               21.98
64               43.91
128              87.99
256              175.36
512              351.74
1024             702.06
2048             1384.25
4096             2149.96
8192             2329.67
16384            980.21
32768            959.38
65536            1232.04
131072           11.79
262144           6.46
524288           2962.36
1048576          2981.89
2097152          2989.36
4194304          2993.12
[sincereli@test04 ~]$
```

- SL (Service Level) to UP mapping

UP = SL & 0x7


- Test RoCE QoS:

```
server$ ib_write_lat -x <GID_IDX> -S <SL> -a
```

```
client$ ib_write_lat -x <GID_IDX> -S <SL> <Server_IP_Addr>
```

```
watch -n 1 'ethtool -S | grep packets | grep prio'
```

# RoCE QoS with Verbs API - Example



```
s6$ ib_write_lat -x 2 -S 2 -a
s5$ ib_write_lat -x 2 -S 2 -a 192.168.100.6
s5$ watch -n 1 'ethtool -S | grep packets | grep prio'
```

```
s5$ ifconfig eth3.100
eth3.100 Link encap:Ethernet HWaddr 00:02:C9:18:97:40
        inet addr:192.168.100.5 Bcast:192.168.100.255 Mask:255.255.255.0

s6$ ifconfig eth3
eth3 Link encap:Ethernet HWaddr 00:02:C9:A0:57:80
        inet addr:192.168.1.6 Bcast:192.168.1.255 Mask:255.255.255.0

s6$ ifconfig eth3.100
eth3.100 Link encap:Ethernet HWaddr 00:02:C9:A0:57:80
        inet addr:192.168.100.6 Bcast:192.168.100.255 Mask:255.255.255.0

s6$ ibv_devinfo -v
hca_id: mlx4_0
transport: InfiniBand (0)
fw_ver: 2.32.5100
raw_eth_odp_caps: NO SUPPORT
max_dct: 0
port: 1
state: PORT_ACTIVE (4)
max_mtu: 4096 (5)
active_width: 4X (2)
active_speed: 10.0 Gbps (4)
phys_state: LINK_UP (5)
GID[ 0]: fe80:0000:0000:0000:0202:c9ff:fea0:5780
GID[ 1]: 0000:0000:0000:0000:0000:ffff:c0a8:0106 -----> 192.168.1.6
GID[ 2]: 0000:0000:0000:0000:0000:ffff:c0a8:6406 -----> 192.168.100.6
```

- Set QoS in user space application

```
rdma_set_option(rdma_cm_id, tos);
```


- TOS\_2\_SL mapping(2 steps)



- Kernel driver handles the mapping

```
skb_prio = rt_tos2priority(id_priv->tos);  
ndev      = ndev->priv_flags & IFF_802_1Q_VLAN ? vlan_dev_real_dev(ndev) : ndev;  
sl        = netdev_get_prio_tc_map(ndev, skb_prio);
```

## ■ Static mapping



```
#define TC_PRIO_BESTEFFORT      0
#define TC_PRIO_FILLER        1
#define TC_PRIO_BULK          2
#define TC_PRIO_INTERACTIVE_BULK 4
#define TC_PRIO_INTERACTIVE   6
#define TC_PRIO_CONTROL       7
#define TC_PRIO_MAX           15

const __u8 ip_tos2prio[16] = {
    TC_PRIO_BESTEFFORT,
    TC_PRIO_BESTEFFORT,
    TC_PRIO_BESTEFFORT,
    TC_PRIO_BESTEFFORT,
    TC_PRIO_BULK,
    TC_PRIO_BULK,
    TC_PRIO_BULK,
    TC_PRIO_BULK,
    TC_PRIO_INTERACTIVE,
    TC_PRIO_INTERACTIVE,
    TC_PRIO_INTERACTIVE,
    TC_PRIO_INTERACTIVE),
    TC_PRIO_INTERACTIVE_BULK,
    TC_PRIO_INTERACTIVE_BULK,
    TC_PRIO_INTERACTIVE_BULK,
    TC_PRIO_INTERACTIVE_BULK
};
```



- Set&Check netdev skprio2up mapping

**Example:**

```
s6$ tc_wrap.py -i eth3 -u 3,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3
```

UP 0

UP 1

UP 2

UP 3

skprio: 0

skprio: 1

skprio: 2 (tos: 8)

skprio: 3

skprio: 4 (tos: 24)

skprio: 5

skprio: 6 (tos: 16)

skprio: 7

skprio: 8

skprio: 9

skprio: 10

skprio: 11

skprio: 12

skprio: 13

skprio: 14

skprio: 15

UP 4

UP 5

UP 6

UP 7

```
s6$ cat /sys/class/net/eth3/qos/skprio2up
```

3 3 3 3 3 3 3 3 3 3 3 3 3 3 3

- [How To Run RoCE and TCP over L2 Enabled with PFC](#)
- RoCE Spec: Annex A16 RoCE.pdf

Q&A

THANK YOU